

```
In [1]: import Pkg
Pkg.activate(@__DIR__)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
using Test
import Convex as cvx
import ECOS
using Random
using ControlSystems
using Plots; plotly()
```

```
Activating environment at `~/Dropbox/My Mac (MacBook Pro (2))/Desktop/CMU/Optimal Control/HW2_S23/Project.toml`
[ Info: Precompiling PlotlyBase [a03496cd-edff-5a9b-9e67-9cda94a718b5]
[ Info: Precompiling PlotlyKaleido [f2990250-8cf9-495f-b13a-cce12b45703c]
└ Warning: backend `PlotlyBase` is not installed.
└ @ Plots ~/.julia/packages/Plots/bMtsB/src/backends.jl:43
└ Warning: backend `PlotlyKaleido` is not installed.
└ @ Plots ~/.julia/packages/Plots/bMtsB/src/backends.jl:43
```

```
Out[1]: Plots.PlotlyBackend()
```

Julia Warnings:

1. For a function `foo(x::Vector)` with 1 input argument, it is not necessary to do `df_dx = FD.jacobian(_x -> foo(_x), x)`. Instead you can just do `df_dx = FD.jacobian(foo, x)`. If you do the first one, it can dramatically slow down your compilation time.
2. Do not define functions inside of other functions like this:

```
function foo(x)
    # main function foo

    function body(x)
        # function inside function (DON'T DO THIS)
        return 2*x
    end

    return body(x)
end
```

This will also slow down your compilation time dramatically.

Q1: Finite-Horizon LQR (50 pts)

For this problem we are going to consider a "double integrator" for our dynamics model. This system has a state $x \in \mathbb{R}^4$, and control $u \in \mathbb{R}^2$, where the state describes the 2D position p and velocity v of an object, and the control is the acceleration a of this object. The state and control are the following:

$$x = [p_1, p_2, v_1, v_2] \quad (1)$$

$$u = [a_1, a_2] \quad (2)$$

And the continuous time dynamics for this system are the following:

$$\dot{x} = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} x + \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 1 & 0 \\ 0 & 1 \end{bmatrix} u \quad (3)u$$

Part A: Discretize the model (5 pts)

Use the matrix exponential (`exp` in Julia) to discretize the continuous time model. See the [first recitation](#) if you're unsure of what to do.

```
In [2]: # double integrator dynamics
function double_integrator_AB(dt)::Tuple{Matrix,Matrix}
    Ac = [0 0 1 0;
           0 0 0 1;
           0 0 0 0;
           0 0 0 0.]
    Bc = [0 0;
           0 0;
           1 0;
           0 1]
    nx, nu = size(Bc)

    # TODO: discretize this linear system using the Matrix Exponential
    matrixExp = exp([Ac*dt Bc*dt; zeros(nu, nx+nu)])
    A = matrixExp[1:nx, 1:nx]
    B = matrixExp[1:nx, (nx+1):end]

    @assert size(A) == (nx,nx)
    @assert size(B) == (nx,nu)

    return A, B
end
```

Out[2]: `double_integrator_AB` (generic function with 1 method)

```
In [3]: @testset "discrete time dynamics" begin
    dt = 0.1
    A,B = double_integrator_AB(dt)

    x = [1,2,3,4.]
    u = [-1,-3.]
    @test isapprox((A*x + B*u), [1.295, 2.385, 2.9, 3.7]; atol = 1e-10)

end
```

	Pass	Total
discrete time dynamics	1	1

Out[3]: `Test.DefaultTestSet("discrete time dynamics", Any[], 1, false, false)`

Part B: Finite Horizon LQR via Convex Optimization (15 pts)

We are now going to solve the finite horizon LQR problem with convex optimization. As we went over in class, this problem requires $Q \in S_+$ (Q is symmetric positive semi-definite) and $R \in S_{++}$ (R is symmetric positive definite). With this, the optimization problem can be stated as the following:

$$\min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \quad (4)$$

$$\text{st } x_1 = x_{\text{IC}} \quad (5)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \quad (6)$$

This problem is a convex optimization problem since the cost function is a convex quadratic and the constraints are all linear equality constraints. We will setup and solve this exact problem using the `Convex.jl` modeling package. (See 2/16 Recitation video for help with this package. [Notebook is here.](#)) Your job in the block below is to fill out a function `Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x_ic)`, where you will form and solve the above optimization problem.

```
In [4]: # utilities for converting to and from vector of vectors <-> matrix
function mat_from_vec(X::Vector{Vector{Float64}})::Matrix
    # convert a vector of vectors to a matrix
    Xm = hcat(X...)
    return Xm
end
function vec_from_mat(Xm::Matrix)::Vector{Vector{Float64}}
    # convert a matrix into a vector of vectors
    X = [Xm[:, i] for i = 1:size(Xm, 2)]
    return X
end
"""
X,U = convex_trajopt(A,B,Q,R,Qf,N,x_ic; verbose = false)
```

This function takes in a dynamics model $x_{k+1} = A*x_k + B*u_k$ and LQR cost Q, R, Q_f , with a horizon size N , and initial condition x_{ic} , and returns the optimal X and U 's from the above optimization problem. You should use the `vec_from_mat` function to convert the solution matrices from cvx into vectors of vectors (`vec_from_mat(X.value)`)

```
function convex_trajopt(A::Matrix,          # A matrix
                        B::Matrix,          # B matrix
                        Q::Matrix,          # cost weight
                        R::Matrix,          # cost weight
                        Qf::Matrix,         # term cost weight
                        N::Int64,           # horizon size
                        x_ic::Vector;       # initial condition
                        verbose = false)
    )::Tuple{Vector{Vector{Float64}}, Vector{Vector{Float64}}}
```

```
# check sizes of everything
nx,nu = size(B)
@assert size(A) == (nx, nx)
@assert size(Q) == (nx, nx)
@assert size(R) == (nu, nu)
@assert size(Qf) == (nx, nx)
@assert length(x_ic) == nx
```

TODO:

```
# create cvx variables where each column is a time step
# hint: x_k = X[:,k], u_k = U[:,k]
X = cvx.Variable(nx, N)
U = cvx.Variable(nu, N - 1)
```

create cost

```

# hint: you can't do x'*Q*x in Convex.jl, you must do cvx.quadform(x,Q)
# hint: add all of your cost terms to `cost`
cost = 0
for k = 1:(N-1)
    xk = X[:, k]
    uk = U[:, k]
    cost += 0.5*cvx.quadform(xk, Q)
    cost += 0.5*cvx.quadform(uk, R)
end

# add terminal cost
cost += 0.5*cvx.quadform(X[:, N], Qf)

# initialize cvx problem
prob = cvx.minimize(cost)

# TODO: initial condition constraint
# hint: you can add constraints to our problem like this:
# prob.constraints += (Gz == h)

#Add initial condition constraint
prob.constraints += X[:,1] == x_ic

#Add dynamics constraints
for k = 1:(N-1)
    prob.constraints += X[:, k+1] == A*X[:, k] + B*U[:, k]
end

# solve problem (silent solver tells us the output)
cvx.solve!(prob, ECOS.Optimizer; silent_solver = !verbose)

if prob.status != cvx.MathOptInterface.OPTIMAL
    error("Convex.jl problem failed to solve for some reason")
end

# convert the solution matrices into vectors of vectors
X = vec_from_mat(X.value)
U = vec_from_mat(U.value)

return X, U
end

```

Out[4]: convex_trajopt

Now let's solve this problem for a given initial condition, and simulate it to see how it does:

In [5]: @testset "LQR via Convex.jl" begin

```

# problem setup stuff
dt = 0.1
tf = 5.0
t_vec = 0:dt:tf
N = length(t_vec)
A,B = double_integrator_AB(dt)
nx,nu = size(B)
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 5*Q

# initial condition
x_ic = [5,7,2,-1.4]

```

```

# setup and solve our convex optimization problem (verbose = true for submission)
Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x_ic; verbose = false)

# TODO: simulate with the dynamics with control Ucvx, storing the
# state in Xsim

# initial condition
Xsim = [zeros(nx) for i = 1:N]
Xsim[1] = 1*x_ic

for k = 1:N-1
    Xsim[k+1] = A*Xsim[k] + B*Ucvx[k]
end

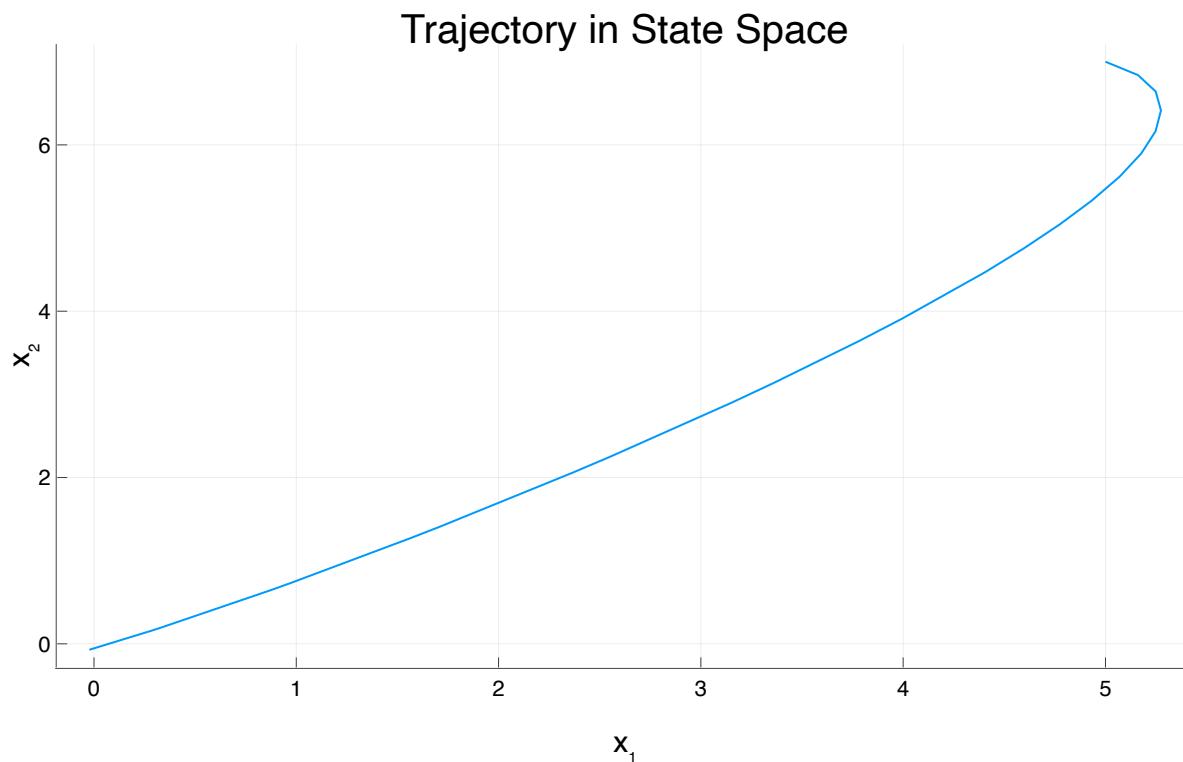
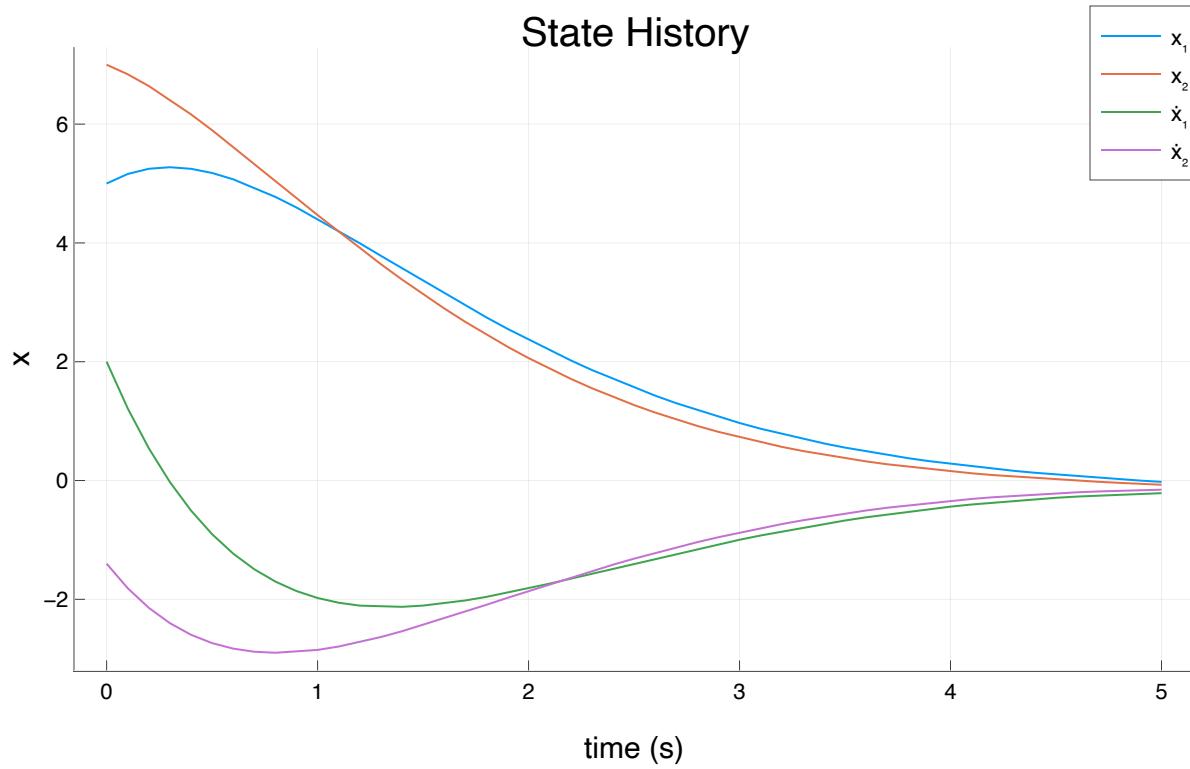
@test length(Xsim) == N
@test norm(Xsim[end])>1e-13
#-----plotting-----
Xsim_m = mat_from_vec(Xsim)

# plot state history
display(plot(t_vec,Xsim_m',label = ["x1" "x2" "dot{x}1" "dot{x}2"],
            title = "State History",
            xlabel = "time (s)", ylabel = "x"))

# plot trajectory in x1 x2 space
display(plot(Xsim_m[1,:],Xsim_m[2,:],
            title = "Trajectory in State Space",
            ylabel = "x2", xlabel = "x1", label = ""))
#-----plotting-----

# tests
@test 1e-14 < maximum(norm.(Xsim .- Xcvx,Inf)) < 1e-3
@test isapprox(Ucvx[1], [-7.8532442316767, -4.127120137234], atol = 1e-3)
@test isapprox(Xcvx[end], [-0.02285990, -0.07140241, -0.21259, -0.1540299], atol = 1e-3)
@test 1e-14 < norm(Xcvx[end] - Xsim[end]) < 1e-3
end

```



Test Summary: | Pass Total
LQR via Convex.jl | 6 6

Out[5]: Test.DefaultTestSet("LQR via Convex.jl", Any[], 6, false, false)

Bellman's Principle of Optimality

Now we will test Bellman's Principle of optimality. This can be phrased in many different ways, but the main gist is that any section of an optimal trajectory must be optimal. Our original optimization problem was the above problem:

$$\min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \quad (7)$$

$$\text{st} \quad x_1 = x_{\text{IC}} \quad (8)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \quad (9)$$

which has a solution $x_{1:N}^*, u_{1:N-1}^*$. Now let's look at optimizing over a subsection of this trajectory. That means that instead of solving for $x_{1:N}, u_{1:N-1}$, we are now solving for $x_{L:N}, u_{L:N-1}$ for some new timestep $1 < L < N$. What we are going to do is take the initial condition from x_L^* from our original optimization problem, and setup a new optimization problem that optimizes over $x_{L:N}, u_{L:N-1}$:

$$\min_{x_{L:N}, u_{L:N-1}} \sum_{i=L}^{N-1} \left[\frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \quad (10)$$

$$\text{st} \quad x_L = x_L^* \quad (11)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = L, L+1, \dots, N-1 \quad (12)$$

In [6]: `@testset "Bellman's Principle of Optimality" begin`

```
# problem setup
dt = 0.1
tf = 5.0
t_vec = 0:dt:tf
N = length(t_vec)
A, B = double_integrator_AB(dt)
nx, nu = size(B)
x0 = [5, 7, 2, -1.4] # initial condition
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 5*Q

# solve for X_{1:N}, U_{1:N-1} with convex optimization
Xcvx1, Ucvx1 = convex_trajopt(A, B, Q, R, Qf, N, x0; verbose = false)

# now let's solve a subsection of this trajectory
L = 18
N_2 = N - L + 1

# here is our updated initial condition from the first problem
x0_2 = Xcvx1[L]
Xcvx2, Ucvx2 = convex_trajopt(A, B, Q, R, Qf, N_2, x0_2; verbose = false)

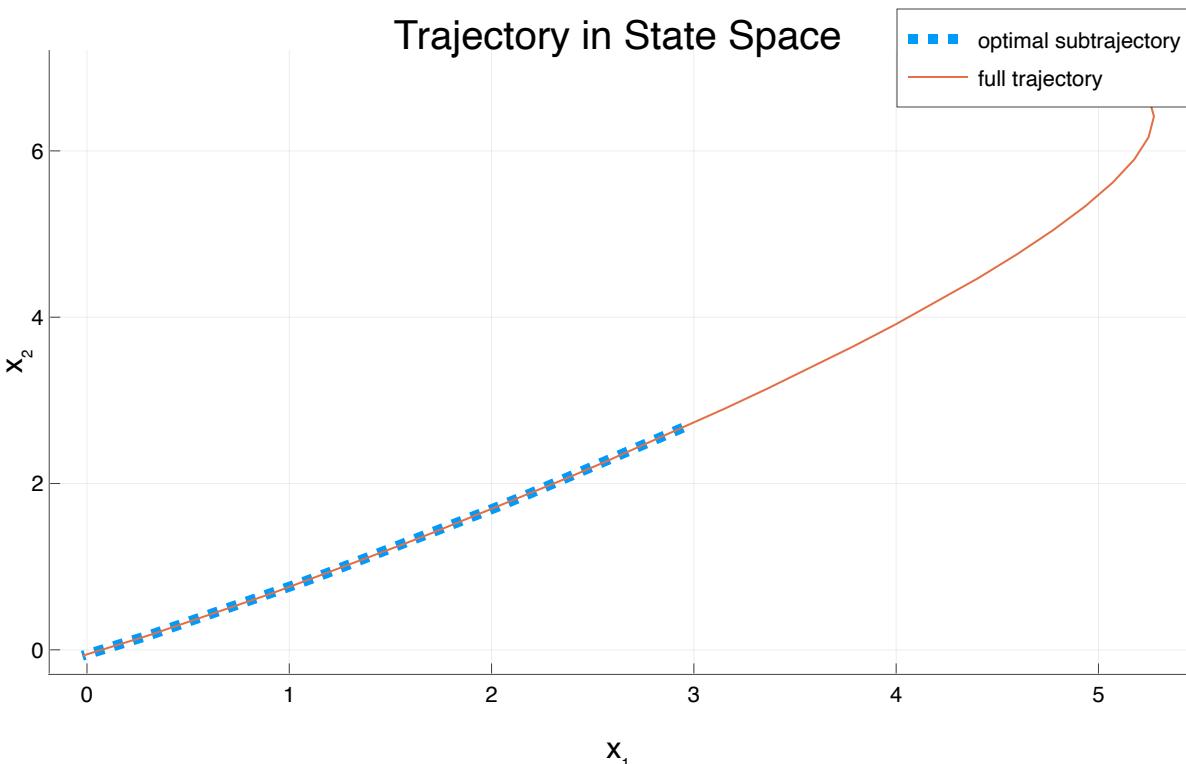
# test if these trajectories match for the times they share
U_error = Ucvx1[L:end] .- Ucvx2
X_error = Xcvx1[L:end] .- Xcvx2
@test 1e-14 < maximum(norm.(U_error)) < 1e-3
@test 1e-14 < maximum(norm.(X_error)) < 1e-3

# -----plotting-----
X1m = mat_from_vec(Xcvx1)
X2m = mat_from_vec(Xcvx2)
plot(X2m[1, :], X2m[2, :], label = "optimal subtrajectory", lw = 5, ls = :dot)
display(plot!(X1m[1, :], X1m[2, :],
            title = "Trajectory in State Space",
            ylabel = "x₂", xlabel = "x₁", label = "full trajectory"))
# -----plotting-----
```

```

@test isapprox(Xcvx1[end], [-0.02285990, -0.07140241, -0.21259, -0.1540299], rtol =
@test 1e-14 < norm(Xcvx1[end] - Xcvx2[end], Inf) < 1e-3
end

```



	Pass	Total
Bellman's Principle of Optimality	4	4

```
Out[6]: Test.DefaultTestSet("Bellman's Principle of Optimality", Any[], 4, false, false)
```

Part C: Finite-Horizon LQR via Riccati (10 pts)

Now we are going to solve the original finite-horizon LQR problem:

$$\min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} x_i^T Q x_i + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} x_N^T Q_f x_N \quad (13)$$

$$\text{st } x_1 = x_{\text{IC}} \quad (14)$$

$$x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \quad (15)$$

with a Riccati recursion instead of convex optimization. We describe our optimal cost-to-go function (aka the Value function) as the following:

$$V_k(x) = \frac{1}{2} x^T P_k x$$

```
In [7]:
```

```
"""
use the Riccati recursion to calculate the cost to go quadratic matrix P and
optimal control gain K at every time step. Return these as a vector of matrices,
where P_k = P[k], and K_k = K[k]
"""

function fhlqr(A::Matrix, # A matrix
                B::Matrix, # B matrix
                Q::Matrix, # cost weight
                R::Matrix, # cost weight
                ...)
```

```

Qf::Matrix, # term cost weight
N::Int64    # horizon size
)::Tuple{Vector{Matrix{Float64}}, Vector{Matrix{Float64}}} # return two m

# check sizes of everything
nx,nu = size(B)
@assert size(A) == (nx, nx)
@assert size(Q) == (nx, nx)
@assert size(R) == (nu, nu)
@assert size(Qf) == (nx, nx)

# instantiate S and K
P = [zeros(nx,nx) for i = 1:N]
K = [zeros(nu,nx) for i = 1:N-1]

# initialize S[N] with Qf
P[N] = deepcopy(Qf)

# Riccati
for i = 1:N-1
    k = N-i
    K[k] = (R+B'*P[k+1]*B)\B'*P[k+1]*A
    P[k] = Q + A'*P[k+1]*(A - B*K[k])
end

return P, K
end

```

Out[7]: fhlqr

In [8]: @testset "Convex trajopt vs LQR" begin

```

# problem stuff
dt = 0.1
tf = 5.0
t_vec = 0:dt:tf
N = length(t_vec)
A,B = double_integrator_AB(dt)
nx,nu = size(B)
x0 = [5,7,2,-1.4] # initial condition
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 5*Q

# solve for X_{1:N}, U_{1:N-1} with convex optimization
Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x0; verbose = false)
P, K = fhlqr(A,B,Q,R,Qf,N)
# now let's simulate using Ucvx
Xsim_cvx = [zeros(nx) for i = 1:N]
Xsim_cvx[1] = 1*x0
Xsim_lqr = [zeros(nx) for i = 1:N]
Xsim_lqr[1] = 1*x0
for i = 1:N-1
    # simulate cvx control
    Xsim_cvx[i+1] = A*Xsim_cvx[i] + B*Ucvx[i]

    # TODO: use your FHLQR control gains K to calculate u_lqr
    # simulate lqr control
    u_lqr = -K[i]*Xsim_lqr[i]
    Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr
end

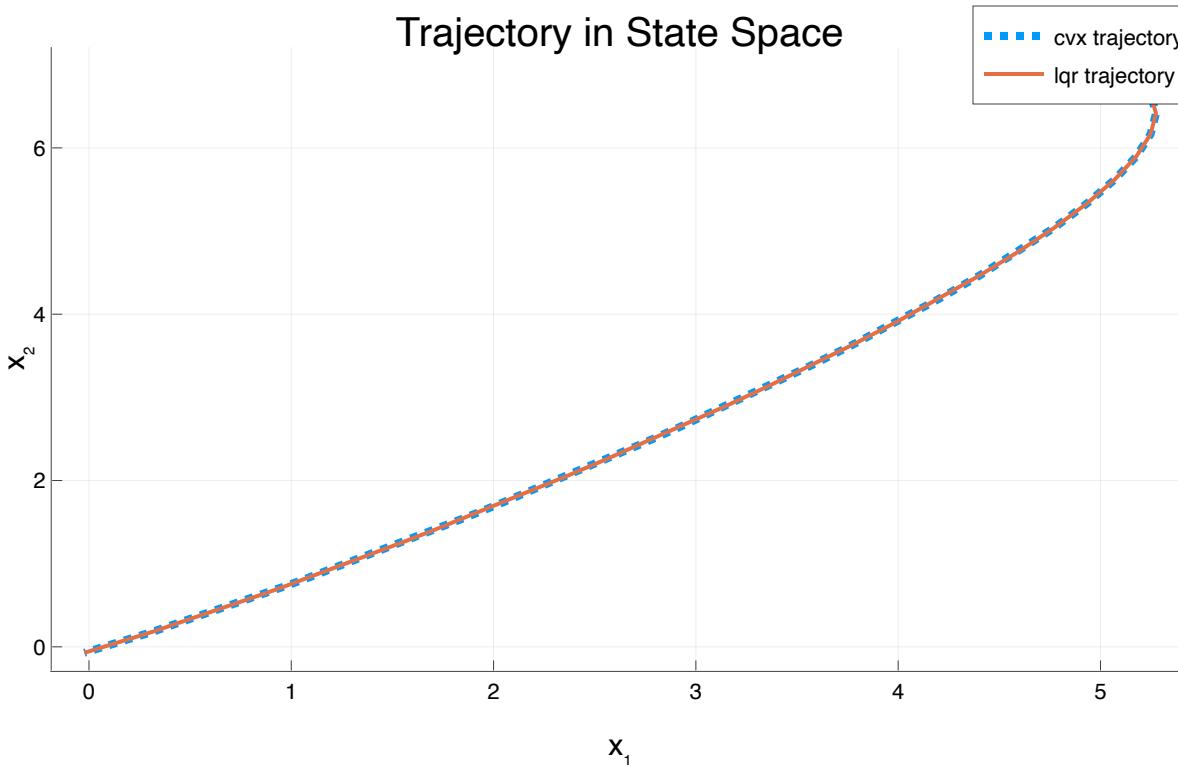
```

```

@test isapprox(Xsim_lqr[end], [-0.02286201, -0.0714058, -0.21259, -0.154030], rtol =
@test 1e-13 < norm(Xsim_lqr[end] - Xsim_cvx[end]) < 1e-3
@test 1e-13 < maximum(norm.(Xsim_lqr - Xsim_cvx)) < 1e-3

# ----- plotting -----
X1m = mat_from_vec(Xsim_cvx)
X2m = mat_from_vec(Xsim_lqr)
# plot trajectory in x1 x2 space
plot(X1m[1,:],X1m[2,:], label = "cvx trajectory", lw = 4, ls = :dot)
display(plot!(X2m[1,:],X2m[2,:],
            title = "Trajectory in State Space",
            ylabel = "x₂", xlabel = "x₁", lw = 2, label = "lqr trajectory"))
# ----- plotting -----
end

```



Test Summary:

	Pass	Total
Convex trajopt vs LQR	3	3

Out[8]: Test.DefaultTestSet("Convex trajopt vs LQR", Any[], 3, false, false)

To emphasize that these two methods for solving the optimization problem result in the same solutions, we are now going to sample initial conditions and run both solutions. You will have to fill in your LQR policy again.

In [9]:

```

import Random
Random.seed!(1)
@testset "Convex trajopt vs LQR" begin

    # problem stuff
    dt = 0.1
    tf = 5.0
    t_vec = 0:dt:tf
    N = length(t_vec)
    A,B = double_integrator_AB(dt)

```

```

nx,nu = size(B)
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 5*Q

plot()
for ic_iter = 1:20
    x0 = [5*randn(2); 1*randn(2)]
    # solve for  $X_{\{1:N\}}$ ,  $U_{\{1:N-1\}}$  with convex optimization
    Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x0; verbose = false)
    P, K = fhlqr(A,B,Q,R,Qf,N)
    Xsim_cvx = [zeros(nx) for i = 1:N]
    Xsim_cvx[1] = 1*x0
    Xsim_lqr = [zeros(nx) for i = 1:N]
    Xsim_lqr[1] = 1*x0
    for i = 1:N-1
        # simulate cvx control
        Xsim_cvx[i+1] = A*Xsim_cvx[i] + B*Ucvx[i]

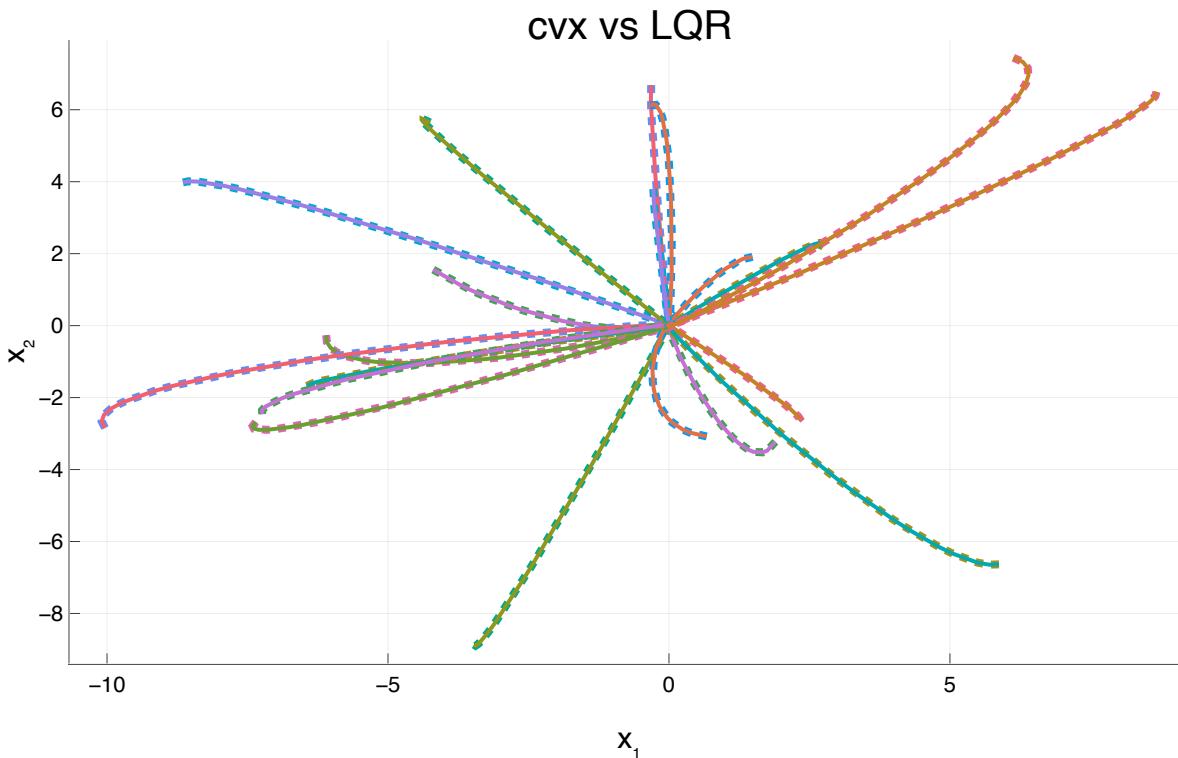
        # TODO: use your FHLQR control gains K to calculate u_lqr
        # simulate lqr control
        u_lqr = -K[i]*Xsim_lqr[i]
        Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr
    end

    @test 1e-13 < norm(Xsim_lqr[end] - Xsim_cvx[end]) < 1e-3
    @test 1e-13 < maximum(norm.(Xsim_lqr - Xsim_cvx)) < 1e-3

    # -----plotting-----
    X1m = mat_from_vec(Xsim_cvx)
    X2m = mat_from_vec(Xsim_lqr)
    plot!(X2m[1,:],X2m[2,:], label = "", lw = 4, ls = :dot)
    plot!(X1m[1,:],X1m[2,:], label = "", lw = 2)
end
display(plot!(title = "cvx vs LQR", ylabel = "x2", xlabel = "x1"))

end

```



	Pass	Total
Convex trajopt vs LQR	40	40

Out[9]: Test.DefaultTestSet("Convex trajopt vs LQR", Any[], 40, false, false)

Part D: Why LQR is so great (10 pts)

Now we are going to emphasize two reasons why the feedback policy from LQR is so useful:

1. It is robust to noise and model uncertainty (the Convex approach would require re-solving of the problem every time the new state differs from the expected state (this is MPC, more on this in Q3))
2. We can drive to any achievable goal state with $u = -K(x - x_{goal})$

First we are going to look at a simulation with the following white noise:

$$x_{k+1} = Ax_k + Bu_k + \text{noise}$$

Where noise $\sim \mathcal{N}(0, \Sigma)$.

In [10]: @testset "Why LQR is great reason 1" begin

```
# problem stuff
dt = 0.1
tf = 7.0
t_vec = 0:dt:tf
N = length(t_vec)
A,B = double_integrator_AB(dt)
nx,nu = size(B)
x0 = [5,7,2,-1.4] # initial condition
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 10*Q

# solve for X_{1:N}, U_{1:N-1} with convex optimization
Xcvx,Ucvx = convex_trajopt(A,B,Q,R,Qf,N,x0; verbose = false)
P, K = fhlqr(A,B,Q,R,Qf,N)
```

```

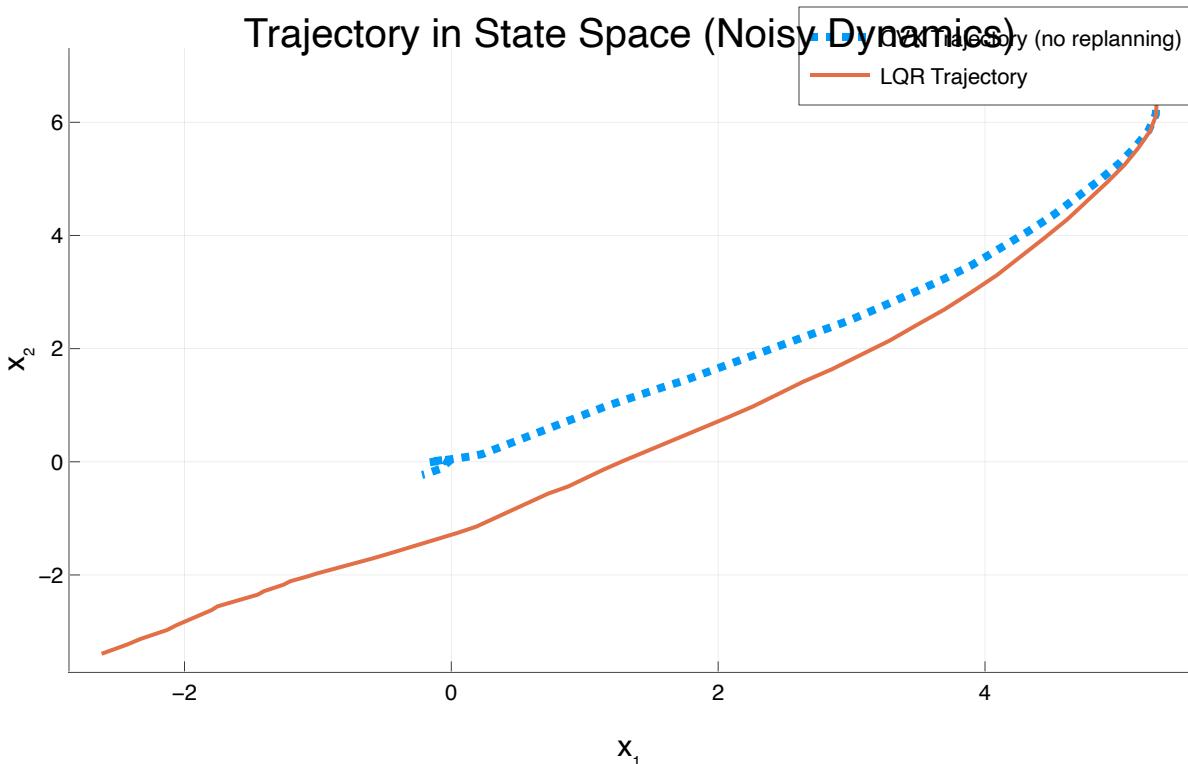
# now let's simulate using Ucvx
Xsim_cvx = [zeros(nx) for i = 1:N]
Xsim_cvx[1] = 1*x0
Xsim_lqr = [zeros(nx) for i = 1:N]
Xsim_lqr[1] = 1*x0
for i = 1:N-1
    # sampled noise to be added after each step
    noise = [.005*randn(2);.1*randn(2)]

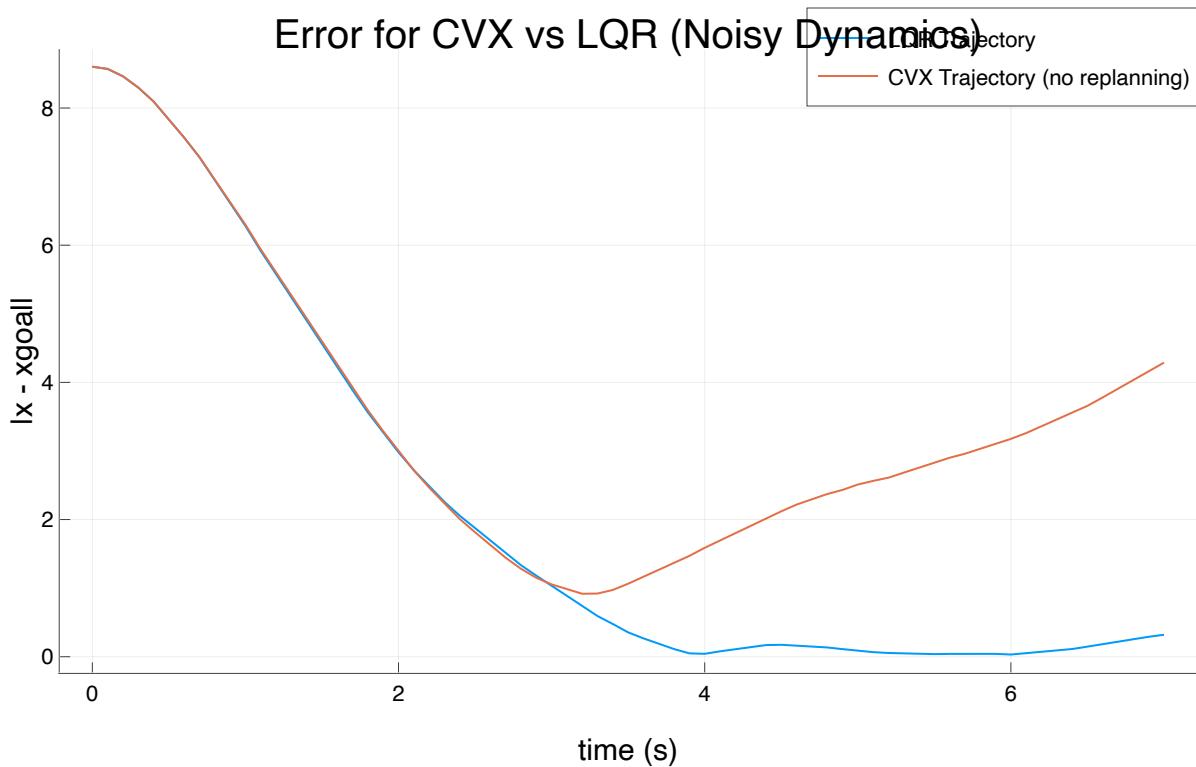
    # simulate cvx control
    Xsim_cvx[i+1] = A*Xsim_cvx[i] + B*Ucvx[i] + noise

    # TODO: use your FHLQR control gains K to calculate u_lqr
    # simulate lqr control
    u_lqr = -K[i]*Xsim_lqr[i]
    Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr + noise
end

# -----plotting-----
X1m = mat_from_vec(Xsim_cvx)
X2m = mat_from_vec(Xsim_lqr)
# plot trajectory in x1 x2 space
plot(X2m[1,:],X2m[2,:], label = "CVX Trajectory (no replanning)", lw = 4, ls = :dot)
display(plot!(X1m[1,:],X1m[2,:],
            title = "Trajectory in State Space (Noisy Dynamics)",
            ylabel = "x2", xlabel = "x1", lw = 2, label = "LQR Trajectory"))
ecvx = [norm(x[1:2]) for x in Xsim_cvx]
elqr = [norm(x[1:2]) for x in Xsim_lqr]
plot(t_vec, elqr, label = "LQR Trajectory", ylabel = "|x - xgoal|",
      xlabel = "time (s)", title = "Error for CVX vs LQR (Noisy Dynamics)")
display(plot!(t_vec, ecvx, label = "CVX Trajectory (no replanning)"))
# -----plotting-----
end

```




Test Summary:

 Why LQR is great reason 1 | [No tests](#)

```
Out[10]: Test.DefaultTestSet("Why LQR is great reason 1", Any[], 0, false, false)
```

```
In [11]: @testset "Why LQR is great reason 2" begin
```

```
# problem stuff
dt = 0.1
tf = 20.0
t_vec = 0:dt:tf
N = length(t_vec)
A, B = double_integrator_AB(dt)
nx, nu = size(B)
x0 = [5, 7, 2, -1.4] # initial condition
Q = diagm(ones(nx))
R = diagm(ones(nu))
Qf = 10*Q

P, K = fhlqr(A, B, Q, R, Qf, N)

# TODO: specify a goal state with 0 velocity within a 5m radius of 0
xgoal = [-3.5, -3.5, 0, 0]
@test norm(xgoal[1:2]) < 5
@test norm(xgoal[3:4]) < 1e-13 # ensure 0 velocity

Xsim_lqr = [zeros(nx) for i = 1:N]
Xsim_lqr[1] = 1*x0

for i = 1:N-1
    # TODO: use your FHLQR control gains K to calculate u_lqr
    # simulate lqr control
    u_lqr = -K[i]*(Xsim_lqr[i] - xgoal)
    Xsim_lqr[i+1] = A*Xsim_lqr[i] + B*u_lqr
end

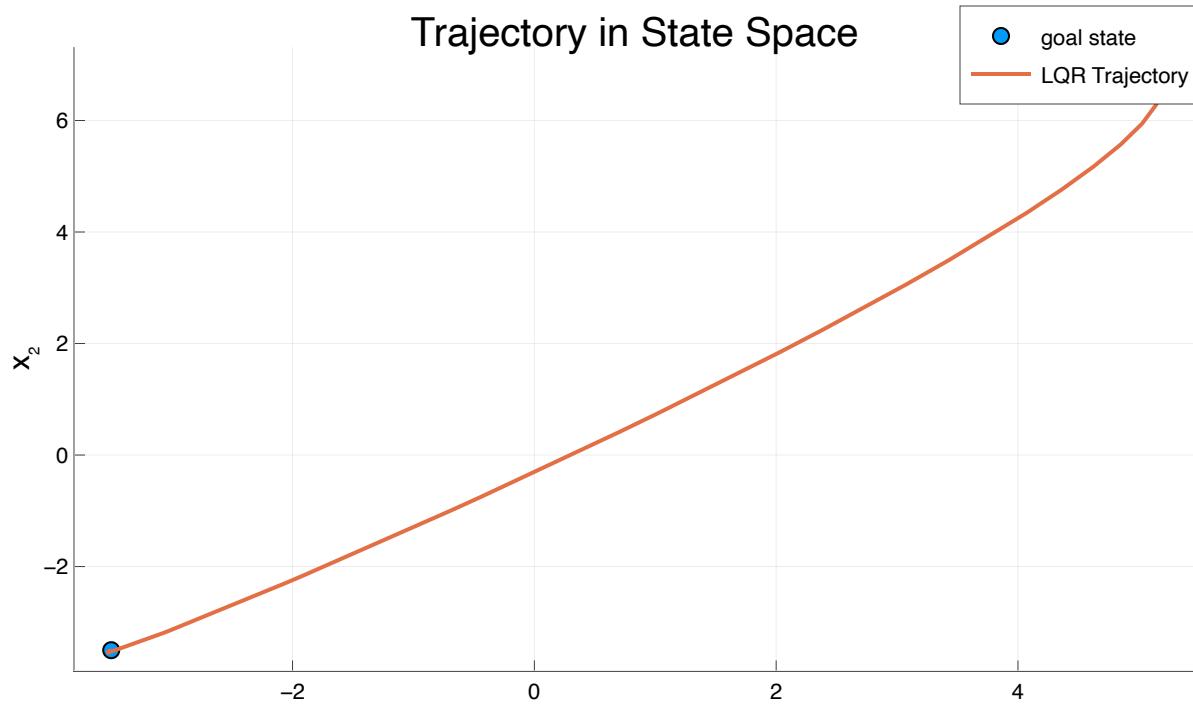
@test norm(Xsim_lqr[end][1:2] - xgoal[1:2]) < .1
```

```

# ----- plotting -----
Xm = mat_from_vec(Xsim_lqr)
plot(xgoal[1:1],xgoal[2:2],seriestype = :scatter, label = "goal state")
display(plot!(Xm[1,:],Xm[2,:],
            title = "Trajectory in State Space",
            ylabel = "x₂", xlabel = "x₁", lw = 2, label = "LQR Trajectory"))

end

```



Test Summary:

	Pass	Total
Why LQR is great reason 2	3	3

Out[11]: Test.DefaultTestSet("Why LQR is great reason 2", Any[], 3, false, false)

Part E: Infinite-horizon LQR (10 pts)

Up until this point, we have looked at finite-horizon LQR which only considers a finite number of timesteps in our trajectory. When this problem is solved with a Riccati recursion, there is a new feedback gain matrix K_k for each timestep. As the length of the trajectory increases, the first feedback gain matrix K_1 will begin to converge on what we call the "infinite-horizon LQR gain". This is the value that K_1 converges to as $N \rightarrow \infty$.

Below, we will plot the values of P and K throughout the horizon and observe this convergence.

```

In [12]: # half vectorization of a matrix
function vech(A)
    return A[tril(true(size(A)))]
end
@testset "P and K time analysis" begin

    # problem stuff
    dt = 0.1
    tf = 10.0
    t_vec = 0:dt:tf
    N = length(t_vec)

```

```

A,B = double_integrator_AB(dt)
nx,nu = size(B)

# cost terms
Q = diagm(ones(nx))
R = .5*diagm(ones(nu))
Qf = randn(nx,nx); Qf = Qf'*Qf + I;

P, K = fhlqr(A,B,Q,R,Qf,N)

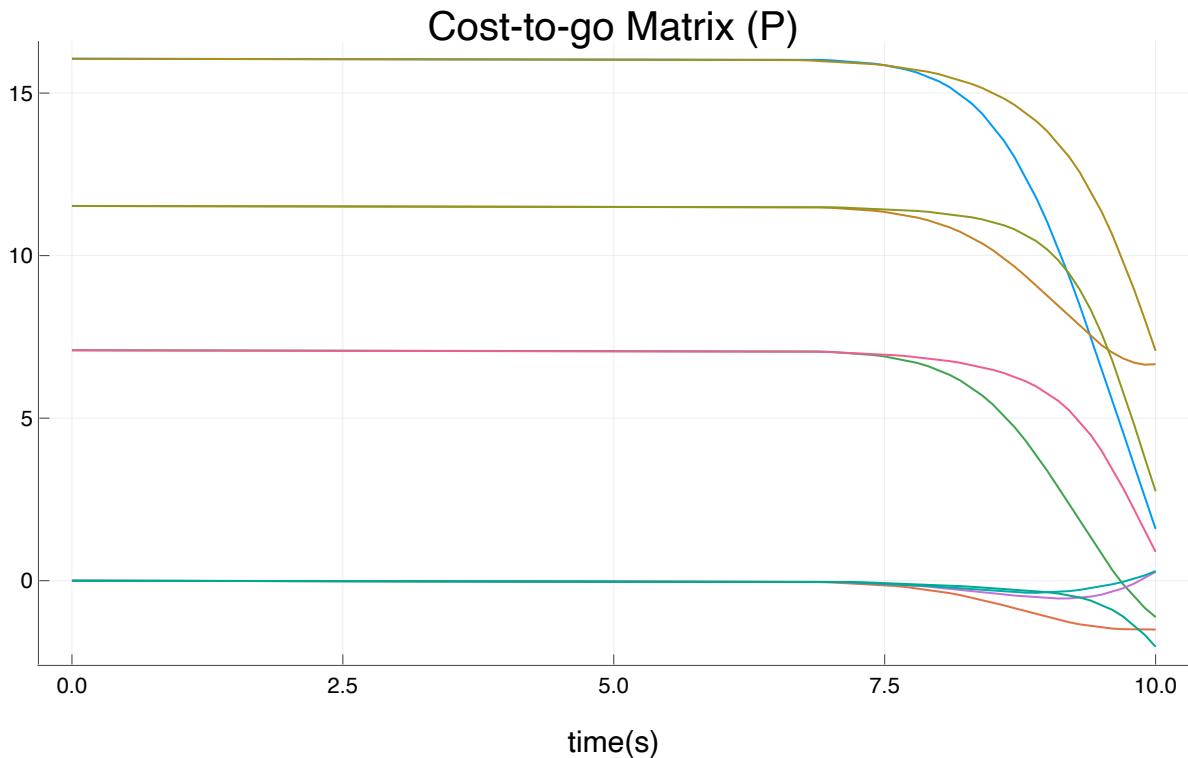
Pm = hcat(vech.(P)...)
Km = hcat(vec.(K)...)

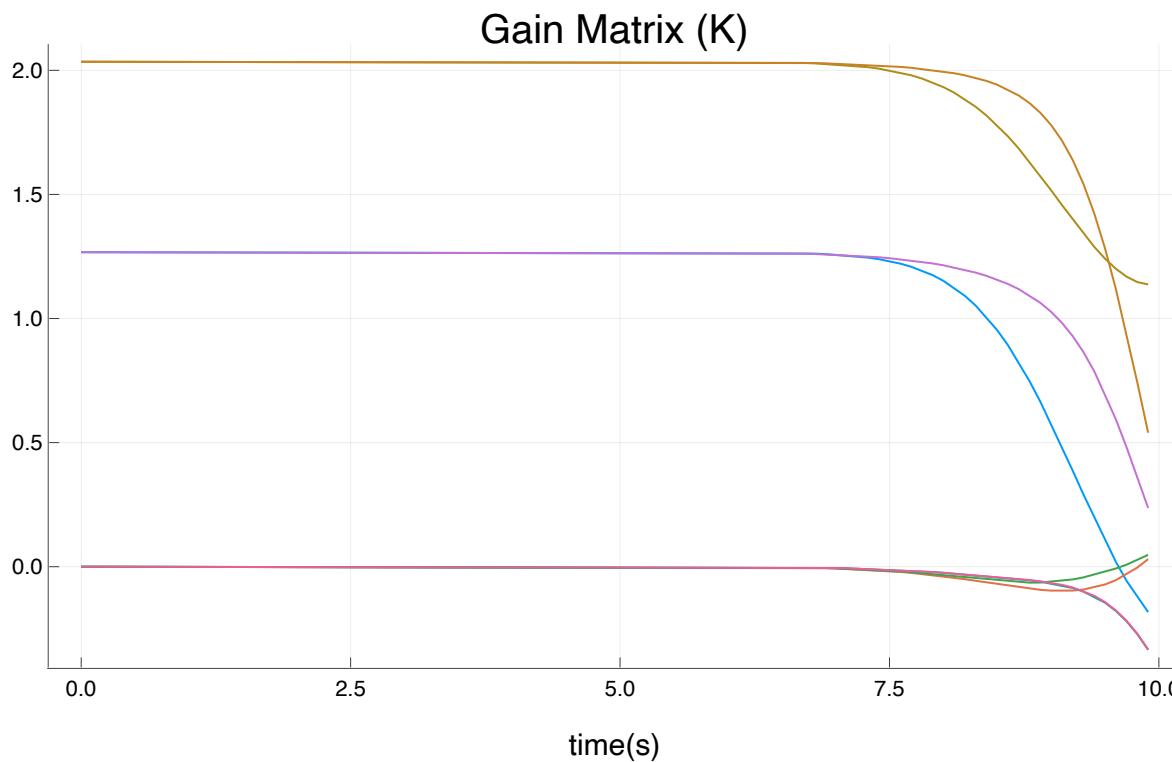
# make sure these things converged
@test 1e-13 < norm(P[1] - P[2]) < 1e-3
@test 1e-13 < norm(K[1] - K[2]) < 1e-3

display(plot(t_vec, Pm', label = "", title = "Cost-to-go Matrix (P)", xlabel = "time"))
display(plot(t_vec[1:end-1], Km', label = "", title = "Gain Matrix (K)", xlabel = "ti

```

end





Test Summary:		Pass	Total
P and K time analysis		2	2

```
Out[12]: Test.DefaultTestSet("P and K time analysis", Any[], 2, false, false)
```

Complete this infinite horizon LQR function where you do a Riccati recursion until the cost to go matrix P converges:

$$\|P_k - P_{k+1}\| \leq \text{tol}$$

And return the steady state P and K .

```
In [14]:
```

```
#####
P,K = ihlqr(A,B,Q,R)

TODO: complete this infinite horizon LQR function where
you do the riccati recursion until the cost to go matrix
P converges to a steady value |P_k - P_{k+1}| <= tol
#####

function ihlqr(A::Matrix,          # vector of A matrices
               B::Matrix,          # vector of B matrices
               Q::Matrix,          # cost matrix Q
               R::Matrix;          # cost matrix R
               max_iter = 1000,    # max iterations for Riccati
               tol = 1e-5           # convergence tolerance
               )::Tuple{Matrix, Matrix} # return two matrices

# get size of x and u from B
nx, nu = size(B)

# initialize S with Q
P = deepcopy(Q)
P_prev = deepcopy(P)

# Riccati
for riccati_iter = 1:max_iter
    K = (R+B'*P*B)\B'*P*A
    P = Q + A'*P*(A - B*K)
```

```

if norm(P - P_prev, 2) < tol
    K
    dlqr(A, B, Q, R)
    return P,K
    break
end
P_prev = P

end
error("ihlqr did not converge")
end
@testset "ihlqr test" begin
    # problem stuff
    dt = 0.1
    A,B = double_integrator_AB(dt)
    nx,nu = size(B)

    # we're just going to modify the system a little bit
    # so the following graphs are still interesting

    Q = diagm(ones(nx))
    R = .5*diagm(ones(nu))
    P, K = ihlqr(A,B,Q,R)

    # check this P is in fact a solution to the Riccati equation
    @test typeof(P) == Matrix{Float64}
    @test typeof(K) == Matrix{Float64}
    @test 1e-13 < norm(Q + K'*R*K + (A - B*K)'P*(A - B*K) - P) < 1e-3

```

end

Test Summary:	Pass	Total
ihlqr test	3	3

Out[14]: Test.DefaultTestSet("ihlqr test", Any[], 3, false, false)

In []:

In [1]:

```

import Pkg
Pkg.activate(@__DIR__)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
import MeshCat as mc
using JLD2
using Test
using Random
using Plots; plotly()
include(joinpath(@__DIR__, "utils/cartpole_animation.jl"))

```

Activating environment at `~/Dropbox/My Mac (MacBook Pro (2))/Desktop/CMU/Optimal Control/HW2_S23/Project.toml`

- ↳ **Warning:** backend `PlotlyBase` is not installed.
- ↳ @ Plots ~/.julia/packages/Plots/bMtsB/src/backends.jl:43
- ↳ **Warning:** backend `PlotlyKaleido` is not installed.
- ↳ @ Plots ~/.julia/packages/Plots/bMtsB/src/backends.jl:43

Out[1]: `animate_cartpole` (generic function with 1 method)

Q2: LQR for nonlinear systems (25 pts)

Linearization warmup

Before we apply LQR to nonlinear systems, we are going to treat our linear system as if it's nonlinear. Specifically, we are going to "approximate" our linear system with a first-order Taylor series, and define a new set of $(\Delta x, \Delta u)$ coordinates. Since our dynamics are linear, this approximation is exact, allowing us to check that we set up the problem correctly.

First, assume our discrete time dynamics are the following:

$$x_{k+1} = f(x_k, u_k)$$

And we are going to linearize about a reference trajectory $\bar{x}_{1:N}, \bar{u}_{1:N-1}$. From here, we can define our delta's accordingly:

$$x_k = \bar{x}_k + \Delta x_k \quad (1)$$

$$u_k = \bar{u}_k + \Delta u_k \quad (2)$$

Next, we are going to approximate our discrete time dynamics function with the following first order Taylor series:

$$x_{k+1} \approx f(\bar{x}_k, \bar{u}_k) + \left[\frac{\partial f}{\partial x} \Big|_{\bar{x}_k, \bar{u}_k} \right] (\bar{x}_k - \bar{x}_k) + \left[\frac{\partial f}{\partial u} \Big|_{\bar{x}_k, \bar{u}_k} \right] (\bar{u}_k - \bar{u}_k)$$

Which we can substitute in our delta notation to get the following:

$$\bar{x}_{k+1} + \Delta x_{k+1} \approx f(\bar{x}_k, \bar{u}_k) + \left[\frac{\partial f}{\partial x} \Big|_{\bar{x}_k, \bar{u}_k} \right] \Delta x_k + \left[\frac{\partial f}{\partial u} \Big|_{\bar{x}_k, \bar{u}_k} \right] \Delta u_k$$

If the trajectory \bar{x}, \bar{u} is dynamically feasible (meaning $\bar{x}_{k+1} = f(\bar{x}_k, \bar{u}_k)$), then we can cancel these equivalent terms on each side of the above equation, resulting in the following:

$$\Delta x_{k+1} \approx \left[\frac{\partial f}{\partial x} \Big|_{x_k, u_k} \right] \Delta x_k + \left[\frac{\partial f}{\partial u} \Big|_{x_k, u_k} \right] \Delta u_k$$

Cartpole

We are now going to look at two different applications of LQR to the nonlinear cartpole system. Given the following description of the cartpole:



(if this image doesn't show up, check out `cartpole.png`)

with a cart position p and pole angle θ . We are first going to linearize the nonlinear discrete dynamics of this system about the point where $p = 0$, and $\theta = 0$ (no velocities), and use an infinite horizon LQR controller about this linearized state to stabilize the cartpole about this goal state. The dynamics of the cartpole are parametrized by the mass of the cart, the mass of the pole, and the length of the pole. To simulate a "sim to real gap", we are going to design our controllers around an estimated set of problem parameters `params_est`, and simulate our system with a different set of problem parameters `params_real`.

In [2]:

```
"""
continuous time dynamics for a cartpole, the state is
x = [p, θ, ̄p, ̄θ]
where p is the horizontal position, and θ is the angle
where θ = 0 has the pole hanging down, and θ = 180 is up.
```

The cartpole is parametrized by a cart mass `mc`, pole mass `mp`, and pole length `l`. These parameters are loaded into a `params::NamedTuple`. We are going to design the controller for a estimated `params_est`, and simulate with `params_real`.

```
"""
function dynamics(params::NamedTuple, x::Vector, u)
    # cartpole ODE, parametrized by params.
```

```
    # cartpole physical parameters
    mc, mp, l = params.mc, params.mp, params.l
    g = 9.81
```

```
    q = x[1:2]
    qd = x[3:4]
```

```
    s = sin(q[2])
    c = cos(q[2])
```

```
    H = [mc+mp mp*l*c; mp*l*c mp*l^2]
    C = [0 -mp*qd[2]*l*s; 0 0]
    G = [0, mp*g*l*s]
    B = [1, 0]
```

```
    qdd = -H\ (C*qd + G - B*u[1])
    return [qd;qdd]
```

```
end
```

```
function rk4(params::NamedTuple, x::Vector, u, dt::Float64)
    # vanilla RK4
```

```

k1 = dt*dynamics(params, x, u)
k2 = dt*dynamics(params, x + k1/2, u)
k3 = dt*dynamics(params, x + k2/2, u)
k4 = dt*dynamics(params, x + k3, u)
x + (1/6)*(k1 + 2*k2 + 2*k3 + k4)
end

```

Out[2]: rk4 (generic function with 1 method)

Part A: Infinite Horizon LQR about an equilibrium (10 pts)

Here we are going to solve for the infinite horizon LQR gain, and use it to stabilize the cartpole about the unstable equilibrium.

```

In [3]: @testset "LQR about eq" begin

    # states and control sizes
    nx = 4
    nu = 1

    # desired x and g (linearize about these)
    xgoal = [0, pi, 0, 0]
    ugoal = [0]

    # initial condition (slightly off of our linearization point)
    x0 = [0, pi, 0, 0] + [1.5, deg2rad(-20), .3, 0]

    # simulation size
    dt = 0.1
    tf = 5.0
    t_vec = 0:dt:tf
    N = length(t_vec)
    X = [zeros(nx) for i = 1:N]
    X[1] = x0

    # estimated parameters (design our controller with these)
    params_est = (mc = 1.0, mp = 0.2, l = 0.5)

    # real parameters (simulate our system with these)
    params_real = (mc = 1.2, mp = 0.16, l = 0.55)

    # TODO: solve for the infinite horizon LQR gain Kinf

    # cost terms
    Q = diagm([1,1,.05,.1])
    Qf = 1*Q
    R = 0.1*diagm(ones(nu))

    #Find linearized A and B matrices
    u0 = zeros(nu)
    A = FD.jacobian(x_ -> dynamics(params_est, x_, ugoal), xgoal)
    B = FD.jacobian(u_ -> dynamics(params_est, xgoal, u_), ugoal)

    P = deepcopy(Qf)
    P_prev = deepcopy(P)
    K = zeros(nx, nu)
    for riccati_iter = 1:1000
        K = (R+B'*P*B)\(B'*P*A)
        P = Q + K'*R*K + (A-B*K)'*P*(A - B*K)
    end

```

```

    if norm(P - P_prev, 2) < 1e-5
        break
    end
    P_prev = P
end

Kinf = K

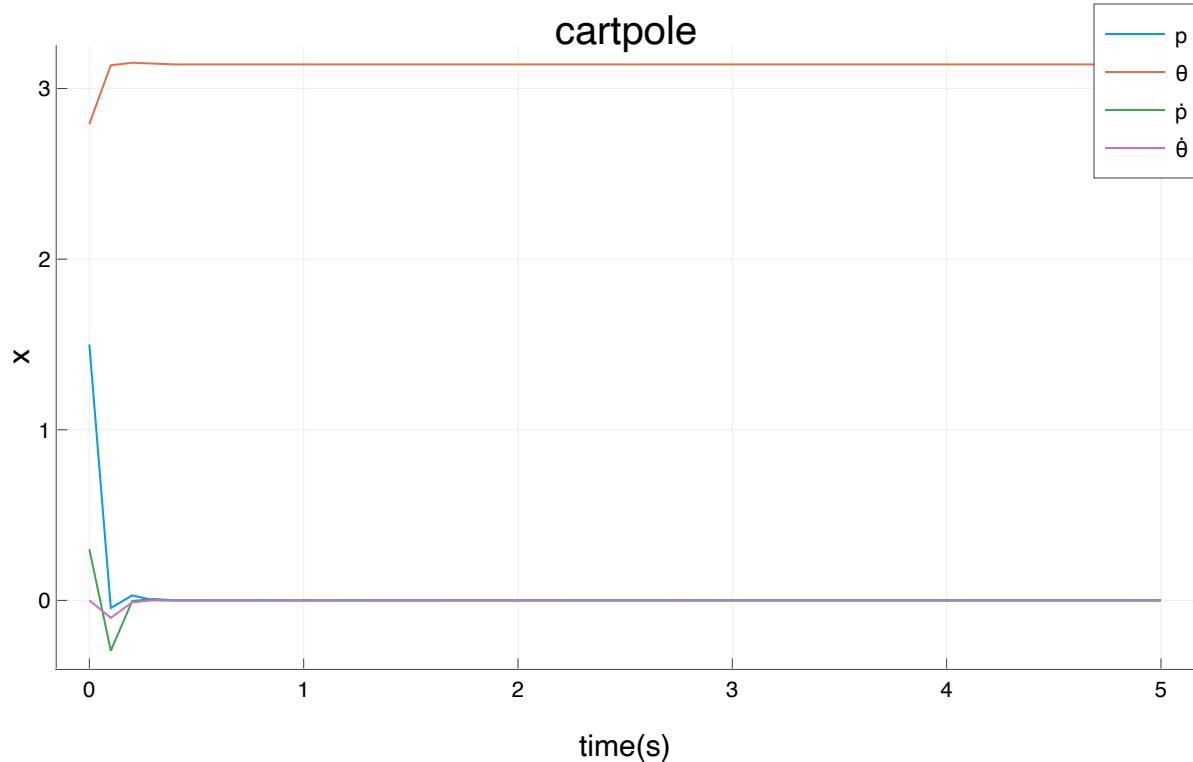
# TODO: simulate this controlled system with rk4(params_real, ...)
for i in 1:N-1
    x_tilde = (X[i] - xgoal)
    uk = ugoal - Kinf*x_tilde
    X[i+1] = X[i] - rk4(params_real, x_tilde, uk, dt)
end

# -----tests and plots/animations-----
@test X[1] == x0
@test norm(X[end])>0
@test norm(X[end] - xgoal) < 0.1

Xm = hcat(X...)
display(plot(t_vec,Xm',title = "cartpole",
            xlabel = "time(s)", ylabel = "x",
            label = ["p" "θ" "ṗ" "θ̇"]))

# animation stuff
display/animate_cartpole(X, dt)
# -----tests and plots/animations-----
end

```



Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:
<http://127.0.0.1:8700>

//

Test Summary:	Pass	Total
LQR about eq	3	3

```
Out[3]: Test.DefaultTestSet("LQR about eq", Any[], 3, false, false)
```

Part B: TVLQR for trajectory tracking (15 pts)

Here we are given a swingup trajectory that works for `params_est`, but will fail to work with `params_real`. To account for this sim to real gap, we are going to track this trajectory with a TVLQR controller.

```
In [4]: @testset "track swingup" begin
```

```
# optimized trajectory we are going to try and track
DATA = load(joinpath(@__DIR__, "swingup.jld2"))
Xbar = DATA["X"]
Ubar = DATA["U"]

# states and controls
nx = 4
nu = 1

# problem size
dt = 0.05
tf = 4.0
t_vec = 0:dt:tf
N = length(t_vec)

# states (initial condition of zeros)
X = [zeros(nx) for i = 1:N]
```

```

X[1] = [0, 0, 0, 0.0]

# make sure we have the same initial condition
@assert norm(X[1] - Xbar[1]) < 1e-12

# real and estimated params
params_est = (mc = 1.0, mp = 0.2, l = 0.5)
params_real = (mc = 1.2, mp = 0.16, l = 0.55)

# TODO: design a time-varying LQR controller to track this trajectory
# use params_est for your control design, and params_real for the simulation

# cost terms
Q = diagm([1,1,.05,.1])
Qf = 10*Q
R = 0.05*diagm(ones(nu))

# TODO: solve for tvlqr gains K

A = [zeros(nx,nx) for i = 1:N-1]
B = [zeros(nx,nu) for i = 1:N-1]

P = [zeros(nx,nx) for i = 1:N]
P[N] = Qf
K = [zeros(nx,nu) for i = 1:N-1]
for k = (N-1):-1:1
    #Linearize system about current trajectory point
    A[k] = FD.jacobian(x_ -> dynamics(params_est, x_, Ubar[k]), Xbar[k])
    B[k] = FD.jacobian(u_ -> dynamics(params_est, Xbar[k], u_), Ubar[k])
    K[k] = (R+B[k]'*P[k+1]*B[k])\((B[k]'*P[k+1]*A[k]))
    P[k] = Q + K[k]'*R*K[k] + (A[k]-B[k]*K[k])'*P[k+1]*(A[k]-B[k]*K[k])
end

# TODO: simulate this controlled system with rk4(params_real, ...)
for i = 1:N-1
    x_tilde = (X[i] - Xbar[i])
    uk = Ubar[i] - K[i]*x_tilde
    X[i+1] = X[i] - rk4(params_real, x_tilde, uk, dt)
end

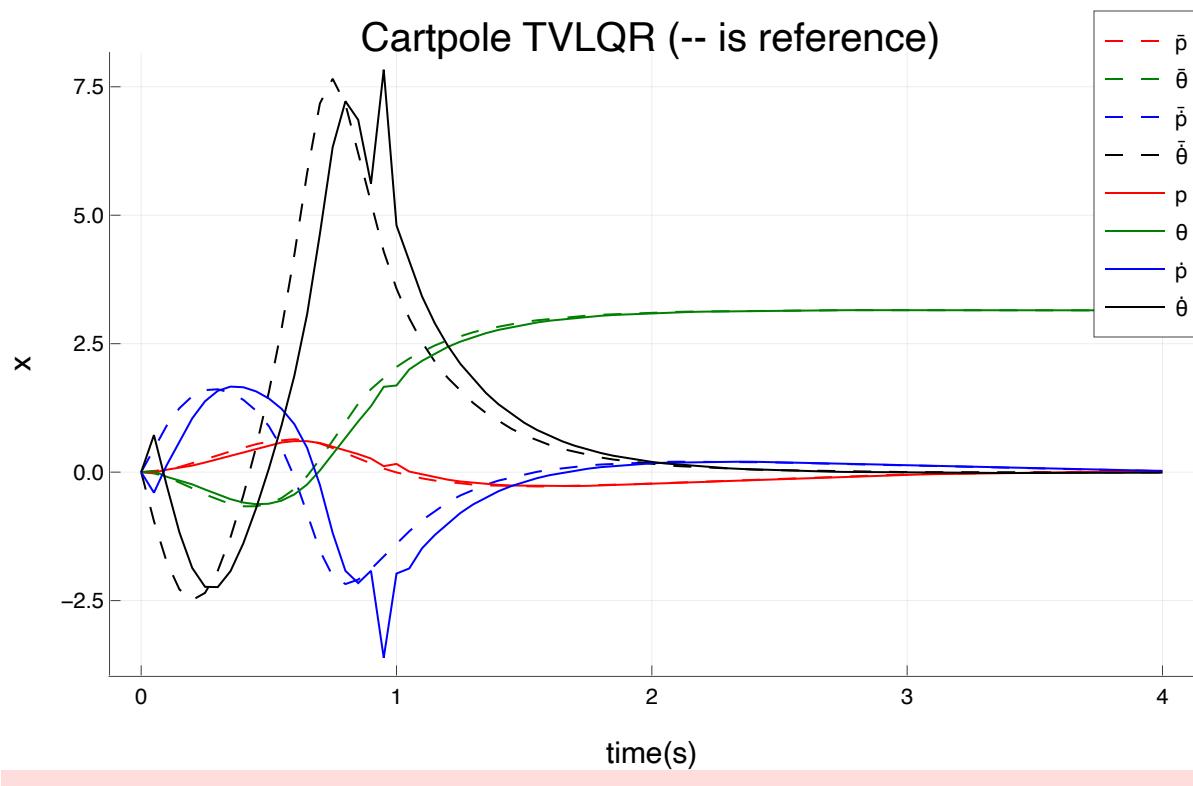
# -----tests and plots/animations-----
xn = X[N]
@test norm(xn)>0
@test 1e-6<norm(xn - Xbar[end])<.2
@test abs(abs(rad2deg(xn[2])) - 180) < 5 # within 5 degrees

Xm = hcat(X...)
Xbarm = hcat(Xbar...)
plot(t_vec,Xbarm',ls=:dash, label = [" $\ddot{\theta}$ " " $\dot{\theta}$ " " $\ddot{x}$ " " $\dot{x}$ "],lc = [:red :green :blue :black]
display(plot!(t_vec,Xm',title = "Cartpole TVLQR (-- is reference)",
            xlabel = "time(s)", ylabel = "x",
            label = ["p" "θ" "dot{p}" "dot{θ}"],lc = [:red :green :blue :black])))

# animation stuff
display/animate_cartpole(X, dt))
# -----tests and plots/animations-----

```

end



Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:
<http://127.0.0.1:8702>

[Open Controls](#)



Out[4]: Test.DefaultTestSet("track swingup", Any[], 3, false, false)

In []:

In [1]:

```
import Pkg
Pkg.activate(@__DIR__)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
import MeshCat as mc
using Test
using Random
import Convex as cvx
import ECOS      # the solver we use in this hw
# import Hypatia # other solvers you can try
# import COSMO   # other solvers you can try
using Plots; plotly()
using ProgressMeter
include(joinpath(@__DIR__, "utils/rendezvous.jl"))
```

```
Activating environment at `~/Dropbox/My Mac (MacBook Pro (2))/Desktop/CMU/Optimal Control/HW2_S23/Project.toml`
└ Warning: backend `PlotlyBase` is not installed.
└ @ Plots ~./julia/packages/Plots/bMtsB/src/backends.jl:43
└ Warning: backend `PlotlyKaleido` is not installed.
└ @ Plots ~./julia/packages/Plots/bMtsB/src/backends.jl:43
```

Out[1]: thruster_model (generic function with 1 method)

Notes:

1. Some of the cells below will have multiple outputs (plots and animations), it can be easier to see everything if you do Cell → All Output → Toggle Scrolling, so that it simply expands the output area to match the size of the outputs.
2. Things in space move very slowly (by design), because of this, you may want to speed up the animations when you're viewing them. You can do this in MeshCat by doing Open Controls → Animations → Time Scale, to modify the time scale. You can also play/pause/scrub from this menu as well.
3. You can move around your view in MeshCat by clicking + dragging, and you can pan with right click + dragging, and zoom with the scroll wheel on your mouse (or trackpad specific alternatives).

In [2]:

```
# utilities for converting to and from vector of vectors <-> matrix
function mat_from_vec(X::Vector{Vector{Float64}})::Matrix
    # convert a vector of vectors to a matrix
    Xm = hcat(X...)
    return Xm
end
function vec_from_mat(Xm::Matrix)::Vector{Vector{Float64}}
    # convert a matrix into a vector of vectors
    X = [Xm[:,i] for i = 1:size(Xm,2)]
    return X
end
```

Out[2]: vec_from_mat (generic function with 1 method)

Is LQR the answer for everything?

Unfortunately, no. LQR is great for problems with true quadratic costs and linear dynamics, but this is a very small subset of convex trajectory optimization problems. While a quadratic cost is common in control, there are other available convex cost functions that may better motivate the desired behavior of the system. These costs can be things like an L1 norm on the control inputs ($\|u\|_1$), or an L2 goal error ($\|x - x_{goal}\|_2$). Also, control problems often have constraints like path constraints, control bounds, or terminal constraints, that can't be handled with LQR. With the addition of these constraints, the trajectory optimization problem is still convex and easy to solve, but we can no longer just get an optimal gain K and apply a feedback policy in these situations.

The solution to this is Model Predictive Control (MPC). In MPC, we are setting up and solving a convex trajectory optimization at every time step, optimizing over some horizon or window into the future, and executing the first control in the solution. To see how this works, we are going to try this for a classic space control problem: the rendezvous.

Q3: Optimal Rendezvous and Docking (50 pts)

In this example, we are going to use convex optimization to control the SpaceX Dragon 1 spacecraft as it docks with the International Space Station (ISS). The dynamics of the Dragon vehicle can be modeled with [Clohessy-Wiltshire equations](#), which is a linear dynamics model in continuous time. The state and control of this system are the following:

$$x = [r_x, r_y, r_z, v_x, v_y, v_z]^T, \quad (1)$$

$$u = [t_x, t_y, t_z]^T, \quad (2)$$

where r is a relative position of the Dragon spacecraft with respect to the ISS, v is the relative velocity, and t is the thrust on the spacecraft. The continuous time dynamics of the vehicle are the following:

$$\dot{x} = \begin{bmatrix} 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 3n^2 & 0 & 0 & 0 & 2n & 0 \\ 0 & 0 & 0 & -2n & 0 & 0 \\ 0 & 0 & -n^2 & 0 & 0 & 0 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} u, \quad (3)$$

where $n = \sqrt{\mu/a^3}$, with μ being the [standard gravitational parameter](#), and a being the semi-major axis of the orbit of the ISS.

We are going to use three different techniques for solving this control problem, the first is LQR, the second is convex trajectory optimization, and the third is convex MPC where we will be able to account for unmodeled dynamics in our system (the "sim to real" gap).

Part A: Discretize the dynamics (5 pts)

Use the matrix exponential to convert the linear ODE into a linear discrete time model (hint: the matrix exponential is just `exp()` in Julia when called on a matrix).

In [3]:

```
function create_dynamics(dt::Real)::Tuple{Matrix,Matrix}
    mu = 3.986004418e14 # standard gravitational parameter
```

```

a = 6971100.0      # semi-major axis of ISS
n = sqrt(mu/a^3)   # mean motion

# continuous time dynamics  $\dot{x} = Ax + Bu$ 
A = [0    0    0    1    0    0;
      0    0    0    0    1    0;
      0    0    0    0    0    1;
      3*n^2 0    0    0    2*n 0;
      0    0    0    -2*n 0    0;
      0    0    -n^2 0    0    0]

B = Matrix([zeros(3,3);0.1*I(3)])

# TODO: convert to discrete time  $X_{k+1} = Ad*x_k + Bd*u_k$ 
nx,nu = size(B)

matrixExp = exp([A*dt B*dt; zeros(nu, nx+nu)])
Ad = matrixExp[1:nx, 1:nx]
Bd = matrixExp[1:nx, (nx+1):end]

return Ad, Bd
end

```

Out[3]: `create_dynamics` (generic function with 1 method)

```

In [4]: @testset "discrete dynamics" begin
    A,B = create_dynamics(1.0)

    x = [1,3,-.3,.2,.4,-.5]
    u = [-.1,.5,.3]

    # test these matrices
    @test isapprox(A*x + B*u, [1.195453, 3.424786, -0.78499972, 0.190925, 0.4495759, -0.
    @test isapprox(det(A), 1, atol = 1e-8)
    @test isapprox(norm(B,Inf), 0.0999999803, atol = 1e-5)

end

Test Summary: | Pass  Total
discrete dynamics |    3      3

```

Out[4]: `Test.DefaultTestSet("discrete dynamics", Any[], 3, false, false)`

Part B: LQR (10 pts)

Now we will take a given reference trajectory `X_ref` and track it with finite-horizon LQR. Remember that finite-horizon LQR is solving this problem:

$$\begin{aligned} \min_{x_{1:N}, u_{1:N-1}} \quad & \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - x_{ref,N})^T Q_f (x_N - x_{ref,N}) \\ \text{st} \quad & x_1 = x_{IC} \\ & x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \end{aligned}$$

where our policy is $u_i = -K_i(x_i - x_{ref,i})$. Use your code from the previous problem with your `fhlqr` function to generate your gain matrices.

One twist we will throw into this is control constraints `u_min` and `u_max`. You should use the function `clamp.(u, u_min, u_max)` to clamp the values of your `u` to be within this range.

If implemented correctly, you should see the Dragon spacecraft dock with the ISS successfully, but only after it crashes through the ISS a little bit.

In [5]:

```
@testset "LQR rendezvous" begin

    # create our discrete time model
    dt = 1.0
    A,B = create_dynamics(dt)

    # get our sizes for state and control
    nx,nu = size(B)

    # initial and goal states
    x0 = [-2;-4;2;0;0;.0]
    xg = [0,-.68,3.05,0,0,0]

    # bounds on U
    u_max = 0.4*ones(3)
    u_min = -u_max

    # problem size and reference trajectory
    N = 120
    t_vec = 0:dt:(N-1)*dt
    X_ref = desired_trajectory_long(x0,xg,200,dt)[1:N]

    # TODO: FHLQR
    Q = diagm(ones(nx))
    R = diagm(ones(nu))
    Qf = 10*Q
    # TODO get K's from fhlqr
    P = [zeros(nx,nx) for i = 1:N]
    K = [zeros(nu,nx) for i = 1:N-1]

    # initialize S[N] with Qf
    P[N] = deepcopy(Qf)

    # Riccati
    for k = (N-1):-1:1
        K[k] = (R+B'*P[k+1]*B)\B'*P[k+1]*A
        P[k] = Q + A'*P[k+1]*(A - B*K[k])
    end

    # simulation
    X_sim = [zeros(nx) for i = 1:N]
    U_sim = [zeros(nu) for i = 1:N-1]
    X_sim[1] = x0
    for i = 1:(N-1)
        # TODO: put LQR control law here
        # make sure to clamp
        x_tilde = X_sim[i] - X_ref[i]
        U_sim[i] = max.(min.(-K[i]*x_tilde, u_max), u_min)

        # simulate 1 step
        X_sim[i+1] = A*X_sim[i] + B*U_sim[i]
    end

    # -----plotting/animation-----
    Xm = mat_from_vec(X_sim)
    Um = mat_from_vec(U_sim)
    display(plot(t_vec,Xm[1:3,:]',title = "Positions (LQR)",
```

```

        xlabel = "time (s)", ylabel = "position (m)",
        label = ["x" "y" "z"]))
display(plot(t_vec,Xm[4:6,:]',title = "Velocities (LQR)",
            xlabel = "time (s)", ylabel = "velocity (m/s)",
            label = ["x" "y" "z"]))
display(plot(t_vec[1:end-1],Um',title = "Control (LQR)",
            xlabel = "time (s)", ylabel = "thrust (N)",
            label = ["x" "y" "z"]))

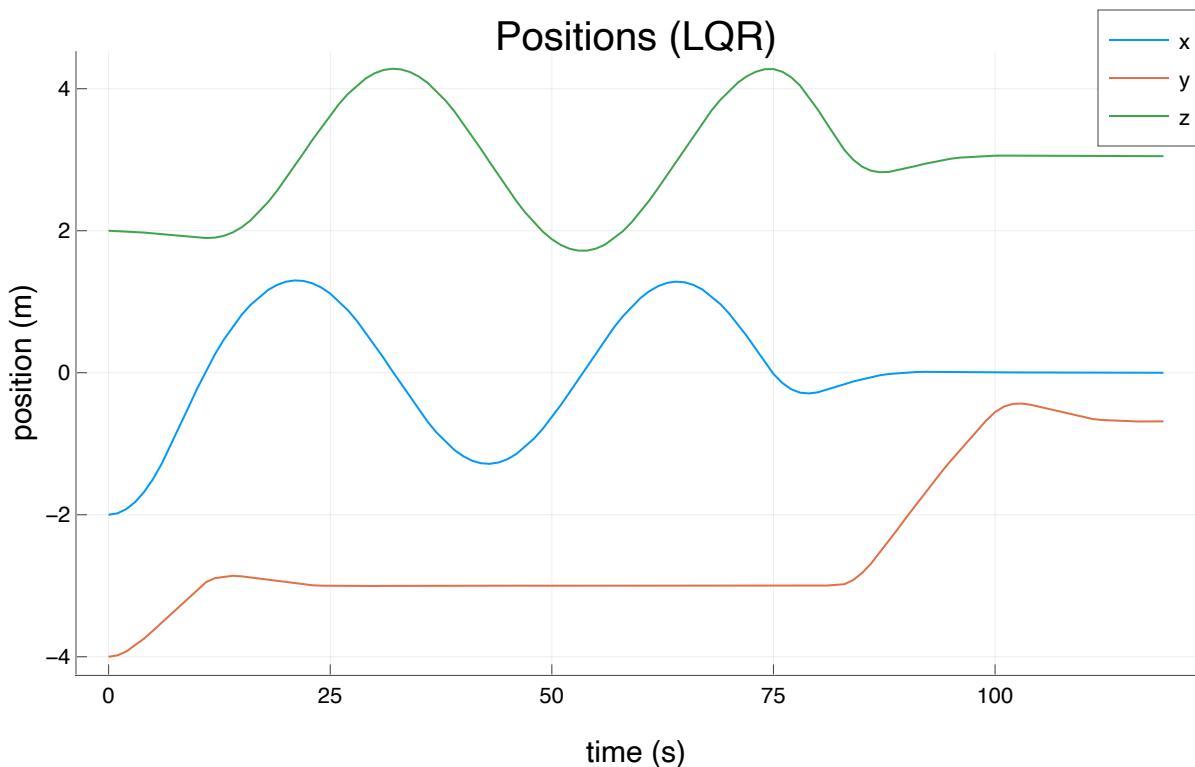
# feel free to toggle `show_reference`
display(animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
# -----plotting/animation-----
```

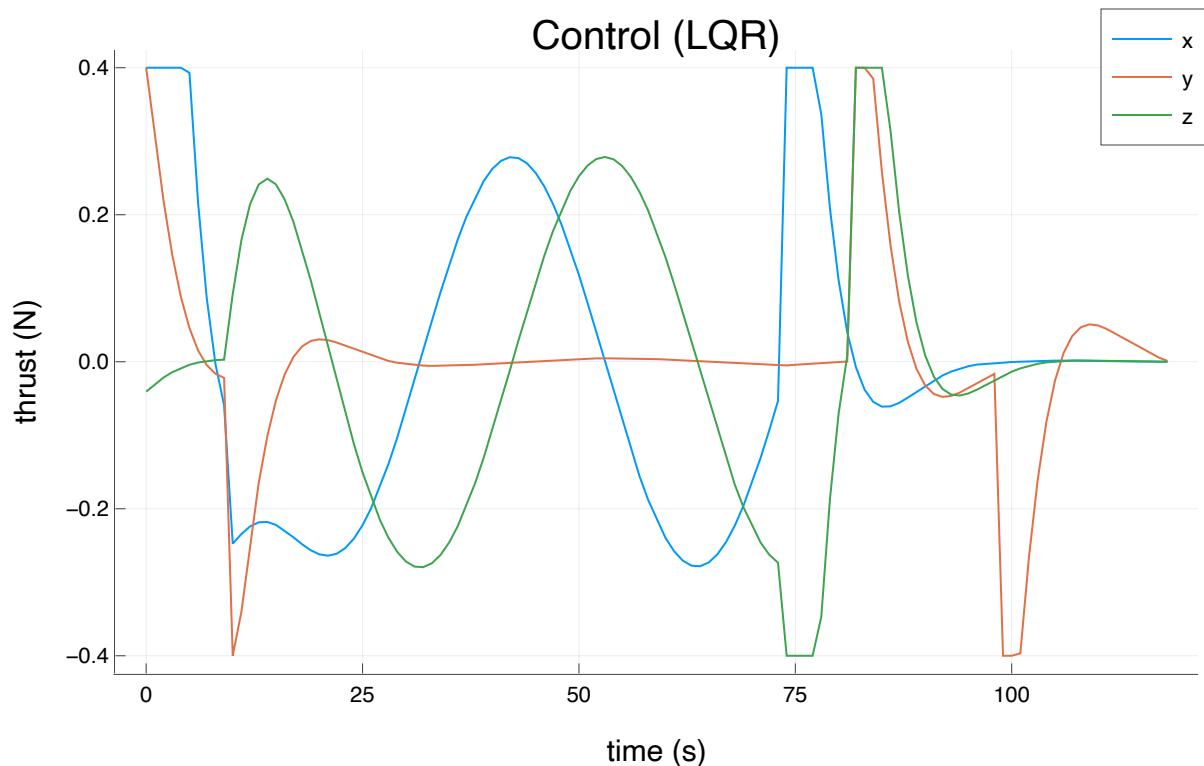
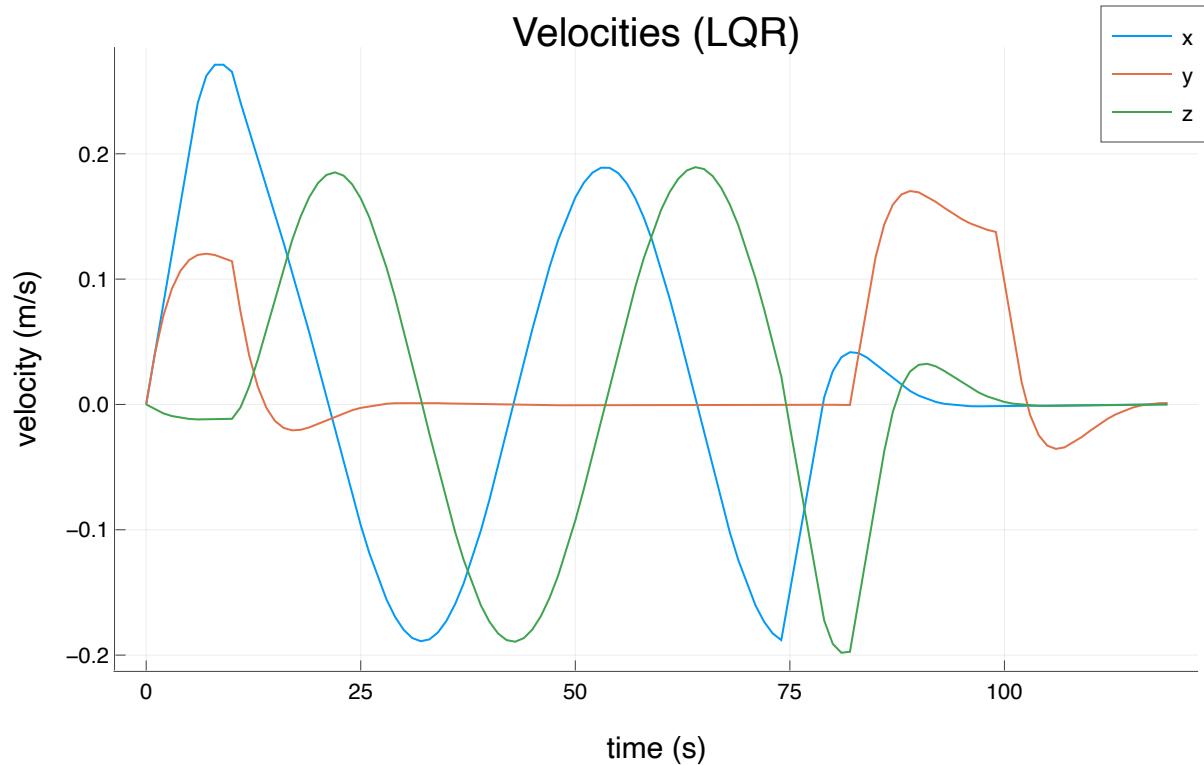
testing

```

xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test norm(X_sim[end] - xg) < .01 # goal
@test (xg[2] + .1) < maximum(ys) < 0 # we should have hit the ISS
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U_sim,Inf)) <= 0.4 # control constraints satisfied
```

end





Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:
http://127.0.0.1:8701

//

Test Summary:	Pass	Total
LQR rendezvous	6	6

```
Out[5]: Test.DefaultTestSet("LQR rendezvous", Any[], 6, false, false)
```

Part C: Convex Trajectory Optimization (15 pts)

Now we are going to assume that we have a perfect model (assume there is no sim to real gap), and that we have a perfect state estimate. With this, we are going to solve our control problem as a convex trajectory optimization problem.

$$\begin{aligned}
 & \min_{x_{1:N}, u_{1:N-1}} \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - x_{ref,i})^T Q (x_i - x_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - x_{ref,N})^T Q_f (x_N - x_{ref,N}) \\
 \text{st } & x_1 = x_{IC} \\
 & x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \\
 & u_{min} \leq u_i \leq u_{max} \quad \text{for } i = 1, 2, \dots, N-1 \\
 & x_i[2] \leq x_{goal}[2] \quad \text{for } i = 1, 2, \dots, N \\
 & x_N = x_{goal}
 \end{aligned}$$

Where we have an LQR cost, an initial condition constraint ($x_1 = x_{IC}$), linear dynamics constraints ($x_{i+1} = Ax_i + Bu_i$), bound constraints on the control ($\leq u_i \leq u_{max}$), an ISS collision constraint ($x_i[2] \leq x_{goal}[2]$), and a terminal constraint ($x_N = x_{goal}$). This problem is convex and we will setup and solve this with `Convex.jl`.

```
In [6]:
```

```
#####
Xcvx,Ucvx = convex_trajopt(A,B,X_ref,x0,xg,u_min,u_max,N)
```

```

setup and solve the above optimization problem, returning
the solutions X and U, after first converting them to
vectors of vectors with vec_from_mat(X.value)
"""
function convex_trajopt(A::Matrix, # discrete dynamics A
                        B::Matrix, # discrete dynamics B
                        X_ref::Vector{Vector{Float64}}, # reference trajectory
                        x0::Vector, # initial condition
                        xg::Vector, # goal state
                        u_min::Vector, # lower bound on u
                        u_max::Vector, # upper bound on u
                        N::Int64, # length of trajectory
                        )::Tuple{Vector{Vector{Float64}}}, Vector{Vector{Float64}}} # ret

# get our sizes for state and control
nx,nu = size(B)

@assert size(A) == (nx, nx)
@assert length(x0) == nx
@assert length(xg) == nx

X_ref = mat_from_vec(X_ref)

# LQR cost
Q = diagm(ones(nx))
R = diagm(ones(nu))

# variables we are solving for
X = cvx.Variable(nx,N)
U = cvx.Variable(nu,N-1)

# TODO: implement cost
obj = 0
for k in 1:(N-1)
    xk_tilde = X[:,k] - X_ref[:,k]
    uk = U[:,k]
    obj += 0.5*cvx.quadform(xk_tilde, Q)
    obj += 0.5*cvx.quadform(uk, R)
end
#Add terminal cost
obj += 0.5*cvx.quadform((X[:,N] - X_ref[:,N]), Q)

# create problem with objective
prob = cvx.minimize(obj)

# TODO: add constraints with prob.constraints +=
prob.constraints += (X[:,1] == x0) #initial condition constraint
for k = 1:(N-1)
    xk = X[:,k]
    uk = U[:,k]
    prob.constraints += (X[:,k+1] == A*xk + B*uk) #dynamics constraints
    #Control Input constraint
    prob.constraints += (uk <= u_max)
    prob.constraints += (uk >= u_min)
    #Position constraint
    prob.constraints += (xk[2] <= xg[2])
end
prob.constraints += (X[2, N] <= xg[2])
prob.constraints += (X[:,N] == xg)

cvx.solve!(prob, ECOS.Optimizer; silent_solver = true)

```

```

X = X.value
U = U.value

Xcvx = vec_from_mat(X)
Ucvx = vec_from_mat(U)

return Xcvx, Ucvx
end

@testset "convex trajopt" begin

# create our discrete time model
dt = 1.0
A,B = create_dynamics(dt)

# get our sizes for state and control
nx,nu = size(B)

# initial and goal states
x0 = [-2;-4;2;0;0;.0]
xg = [0,-.68,3.05,0,0,0]

# bounds on U
u_max = 0.4*ones(3)
u_min = -u_max

# problem size and reference trajectory
N = 100
t_vec = 0:dt:(N-1)*dt
X_ref = desired_trajectory(x0,xg,N,dt)

# solve convex trajectory optimization problem
X_cvx, U_cvx = convex_trajopt(A,B,X_ref, x0,xg,u_min,u_max,N)

X_sim = [zeros(nx) for i = 1:N]
X_sim[1] = x0
for i = 1:N-1
    X_sim[i+1] = A*X_sim[i] + B*U_cvx[i]
end

# -----plotting/animation-----
Xm = mat_from_vec(X_sim)
Um = mat_from_vec(U_cvx)
display(plot(t_vec,Xm[1:3,:]',title = "Positions",
            xlabel = "time (s)", ylabel = "position (m)",
            label = ["x" "y" "z"]))
display(plot(t_vec,Xm[4:6,:]',title = "Velocities",
            xlabel = "time (s)", ylabel = "velocity (m/s)",
            label = ["x" "y" "z"]))
display(plot(t_vec[1:end-1],Um',title = "Control",
            xlabel = "time (s)", ylabel = "thrust (N)",
            label = ["x" "y" "z"]))

display/animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
# -----plotting/animation-----

@test maximum(norm.(X_sim .- X_cvx, Inf)) < 1e-3
@test norm(X_sim[end] - xg) < 1e-3 # goal

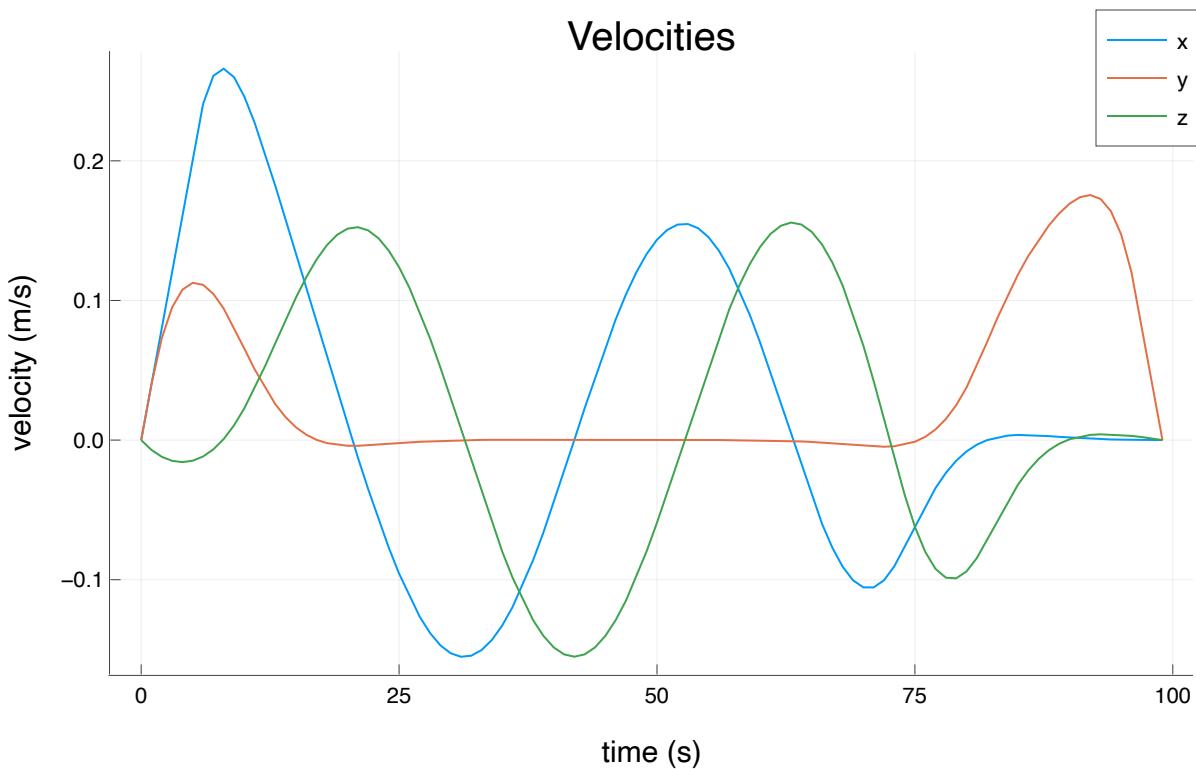
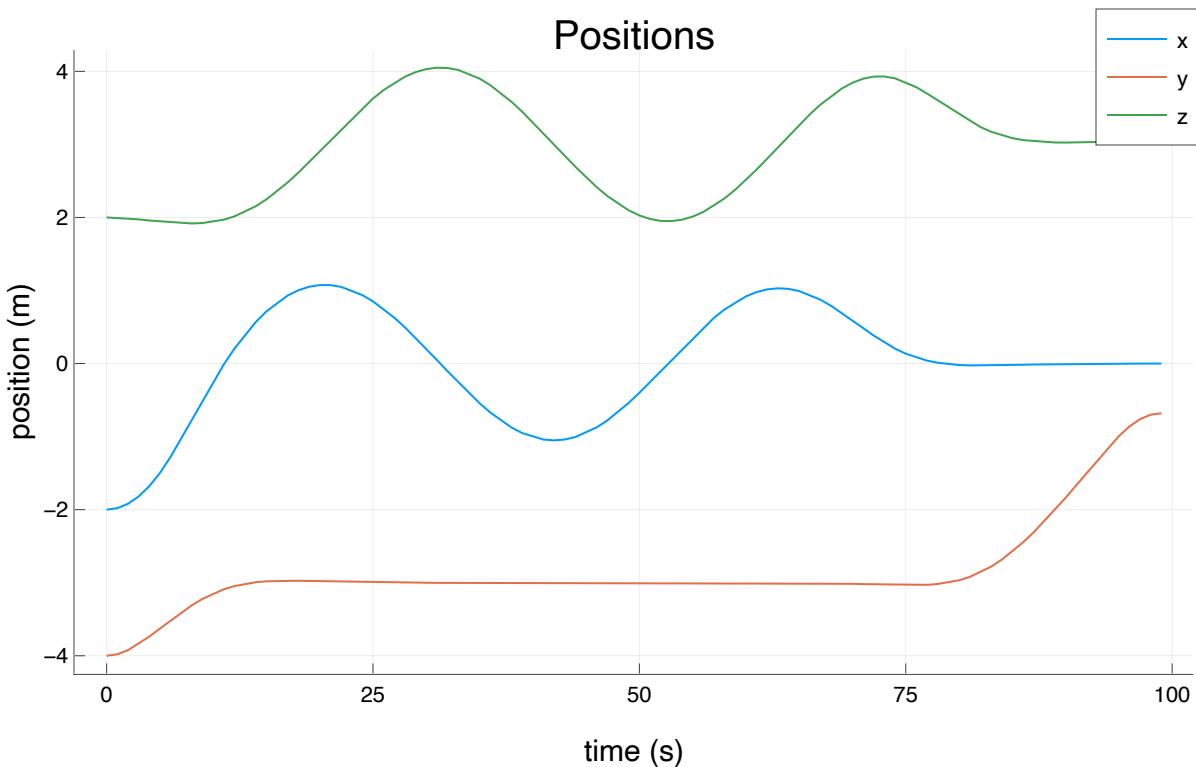
```

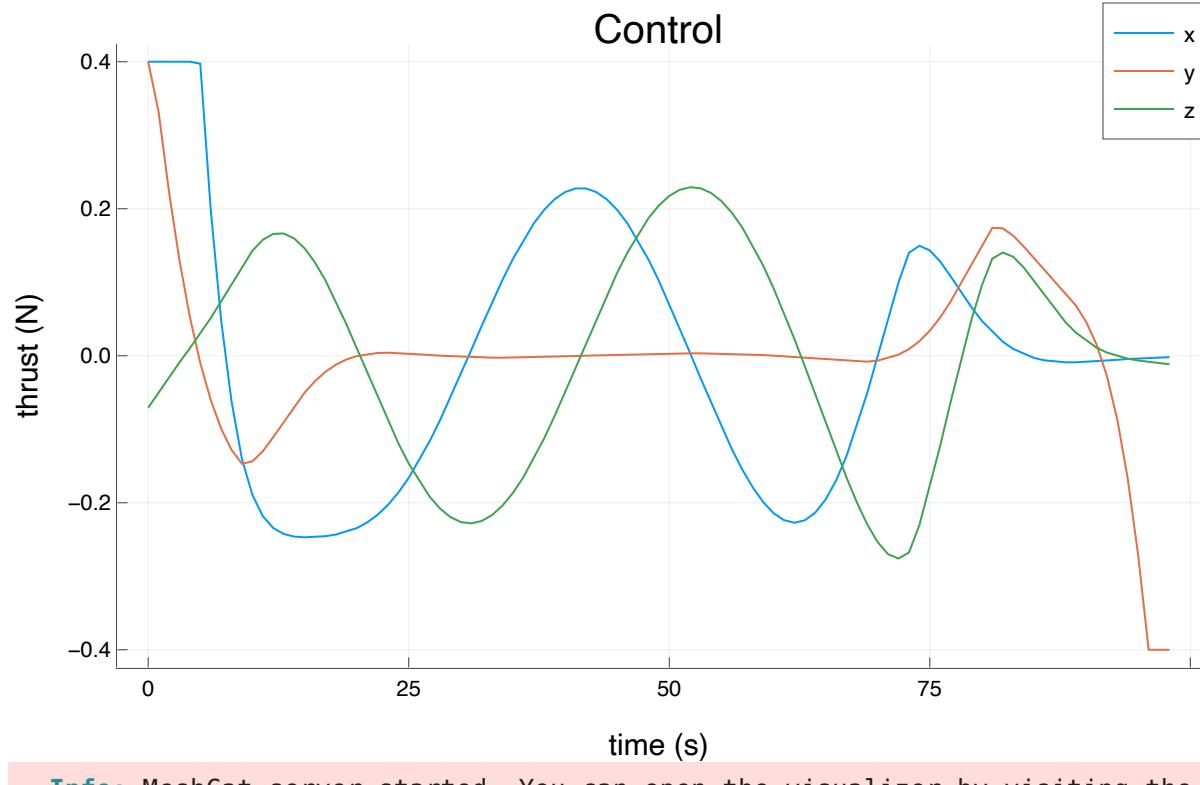
```

xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test maximum(ys) <= (xg[2] + 1e-3)
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U_cvx,Inf)) <= 0.4 + 1e-3 # control constraints satisfied

end

```





Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:
<http://127.0.0.1:8703>

[Open Controls](#)

Test Summary: | Pass Total
convex trajopt | 7 7

Out[6]: `Test.DefaultTestSet("convex trajopt", Any[], 7, false, false)`

Part D: Convex MPC (20 pts)

In part C, we solved for the optimal rendezvous trajectory using convex optimization, and verified it by simulating it in an open loop fashion (no feedback). This was made possible because we assumed that our linear dynamics were exact, and that we had a perfect estimate of our state. In reality, there are many issues that would prevent this open loop policy from being successful, here are a few:

- imperfect state estimation
- unmodeled dynamics
- misalignments
- actuator uncertainties

Together, these factors result in a "sim to real" gap between our simulated model, and the real model. Because there will always be a sim to real gap, we can't just execute open loop policies and expect them to be successful. What we can do, however, is use Model Predictive Control (MPC) that combines some of the ideas of feedback control with convex trajectory optimization.

A convex MPC controller will set up and solve a convex optimization problem at each time step that incorporates the current state estimate as an initial condition. For a trajectory tracking problem like this rendezvous, we want to track x_{ref} , but instead of optimizing over the whole trajectory, we will only consider a sliding window of size N_{mpc} (also called a horizon). If $N_{mpc} = 20$, this means our convex MPC controller is reasoning about the next 20 steps in the trajectory. This optimization problem at every timestep will start by taking the relevant reference trajectory at the current window from the current step i , to the end of the window $i + N_{mpc} - 1$. This slice of the reference trajectory that applies to the current MPC window will be called $\tilde{x}_{ref} = x_{ref}[i, (i + N_{mpc} - 1)]$.

$$\begin{aligned} \min_{x_{1:N}, u_{1:N-1}} \quad & \sum_{i=1}^{N-1} \left[\frac{1}{2} (x_i - \tilde{x}_{ref,i})^T Q (x_i - \tilde{x}_{ref,i}) + \frac{1}{2} u_i^T R u_i \right] + \frac{1}{2} (x_N - \tilde{x}_{ref,N})^T Q_f (x_N - \tilde{x}_{ref,N}) \\ \text{st} \quad & x_1 = x_{IC} \\ & x_{i+1} = Ax_i + Bu_i \quad \text{for } i = 1, 2, \dots, N-1 \\ & u_{min} \leq u_i \leq u_{max} \quad \text{for } i = 1, 2, \dots, N-1 \\ & x_i[2] \leq x_{goal}[2] \quad \text{for } i = 1, 2, \dots, N \end{aligned}$$

where N in this case is N_{mpc} . This allows for the MPC controller to "think" about the future states in a way that the LQR controller cannot. By updating the reference trajectory window (\tilde{x}_{ref}) at each step and updating the initial condition (x_{IC}), the MPC controller is able to "react" and compensate for the sim to real gap.

You will now implement a function `convex_mpc` where you setup and solve this optimization problem at every timestep, and simply return u_1 from the solution.

In [7]:

```
"""
`u = convex_mpc(A,B,X_ref_window,xic,xg,u_min,u_max,N_mpc)`

setup and solve the above optimization problem, returning the
first control u_1 from the solution (should be a length nu
Vector{Float64}).
"""

function convex_mpc(A::Matrix, # discrete dynamics matrix A
```

```

B::Matrix, # discrete dynamics matrix B
X_ref_window::Vector{Vector{Float64}}, # reference trajectory for th
xic::Vector, # current state x
xg::Vector, # goal state
u_min::Vector, # lower bound on u
u_max::Vector, # upper bound on u
N_mpc::Int64, # length of MPC window (horizon)
)::Vector{Float64} # return the first control command of the solved

```

```

# get our sizes for state and control
nx,nu = size(B)

```

```

# check sizes
@assert size(A) == (nx, nx)
@assert length(xic) == nx
@assert length(xg) == nx
@assert length(X_ref_window) == N_mpc

```

```
X_ref_window = mat_from_vec(X_ref_window)
```

```
# LQR cost
```

```
Q = diagm(ones(nx))
R = diagm(ones(nu))
```

```
# variables we are solving for
X = cvx.Variable(nx,N_mpc)
U = cvx.Variable(nu,N_mpc-1)
```

```
# TODO: implement cost
```

```
obj = 0
for k in 1:(N_mpc-1)
    xk_tilde = X[:,k] - X_ref_window[:,k]
    uk = U[:,k]
    obj += 0.5*cvx.quadform(xk_tilde, Q)
    obj += 0.5*cvx.quadform(uk, R)
end
```

```
#Add terminal cost
```

```
obj += 0.5*cvx.quadform((X[:,N_mpc] - X_ref_window[:,N_mpc]), Q)
```

```
# create problem with objective
```

```
prob = cvx.minimize(obj)
```

```
# TODO: add constraints with prob.constraints +=
```

```
prob.constraints += (X[:,1] == xic) #initial condition constraint
```

```
for k = 1:(N_mpc-1)
    xk = X[:,k]
    uk = U[:,k]
    prob.constraints += (X[:,k+1] == A*xk + B*uk) #dynamics constraints
    #Control Input constraint
    prob.constraints += (uk <= u_max)
    prob.constraints += (uk >= u_min)
    #Position constraint
    prob.constraints += (xk[2] <= xg[2])
end
```

```
prob.constraints += (X[2,N_mpc] <= xg[2])
prob.constraints += (X[:,N_mpc] == xg)
```

```
# solve problem
```

```
cvx.solve!(prob, ECOS.Optimizer; silent_solver = true)
```

```
# get X and U solutions
```

```

X = X.value
U = U.value

# return first control U
return U[:,1]
end

@testset "convex mpc" begin

# create our discrete time model
dt = 1.0
A,B = create_dynamics(dt)

# get our sizes for state and control
nx,nu = size(B)

# initial and goal states
x0 = [-2;-4;2;0;0;.0]
xg = [0,-.68,3.05,0,0,0]

# bounds on U
u_max = 0.4*ones(3)
u_min = -u_max

# problem size and reference trajectory
N = 100
t_vec = 0:dt:(N-1)*dt
X_ref = [desired_trajectory(x0,xg,N,dt)..., [xg for i = 1:N]...]
@show size(X_ref)

# MPC window size
N_mpc = 20

# sim size and setup
N_sim = N + 20
t_vec = 0:dt:(N_sim-1)*dt
X_sim = [zeros(nx) for i = 1:N_sim]
X_sim[1] = x0
U_sim = [zeros(nu) for i = 1:N_sim-1]

# simulate
@showprogress "simulating" for i = 1:N_sim-1

    # get state estimate
    xi_estimate = state_estimate(X_sim[i], xg)

    # TODO: given a window of N_mpc timesteps, get current reference trajectory
    X_ref_tilde = X_ref[i:i+N_mpc-1]

    # TODO: call convex mpc controller with state estimate
    u_mpc = convex_mpc(A, B, X_ref_tilde, xi_estimate, X_ref_tilde[end], u_min, u_max)

    # commanded control goes into thruster model where it gets modified
    U_sim[i] = thruster_model(X_sim[i], xg, u_mpc)

    # simulate one step
    X_sim[i+1] = A*X_sim[i] + B*U_sim[i]
end

# -----plotting/animation-----

```

```

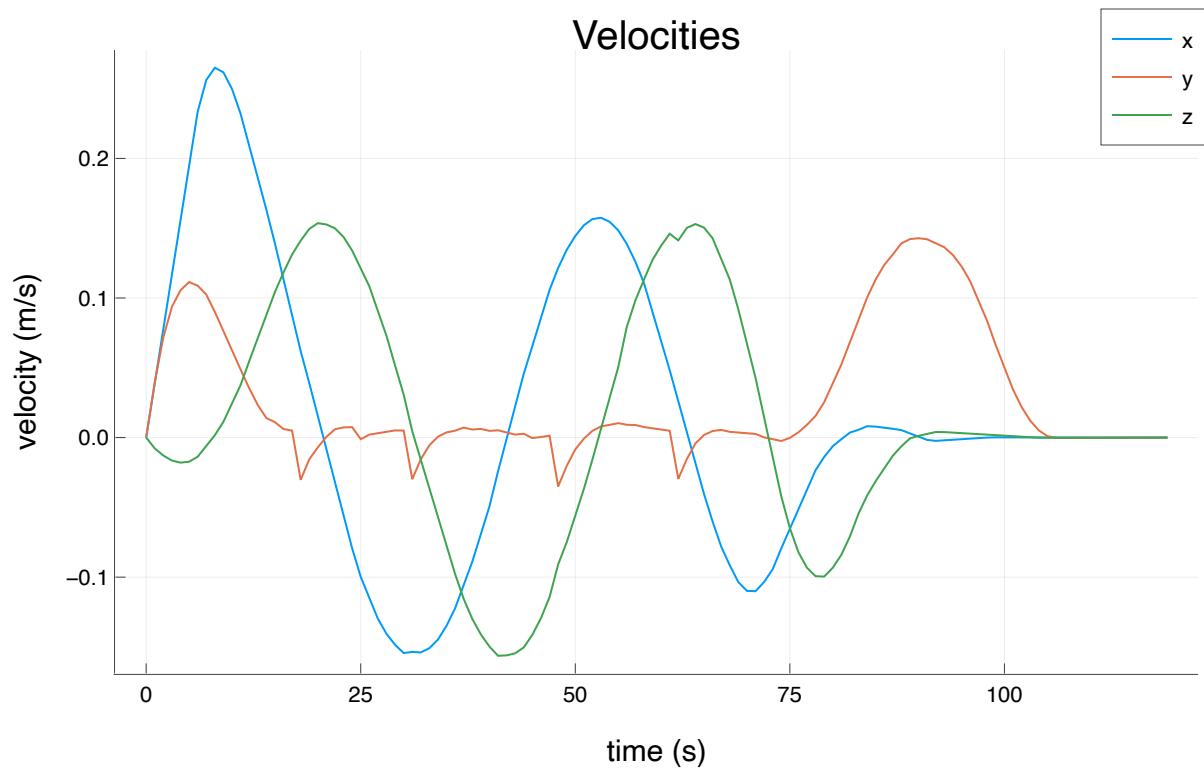
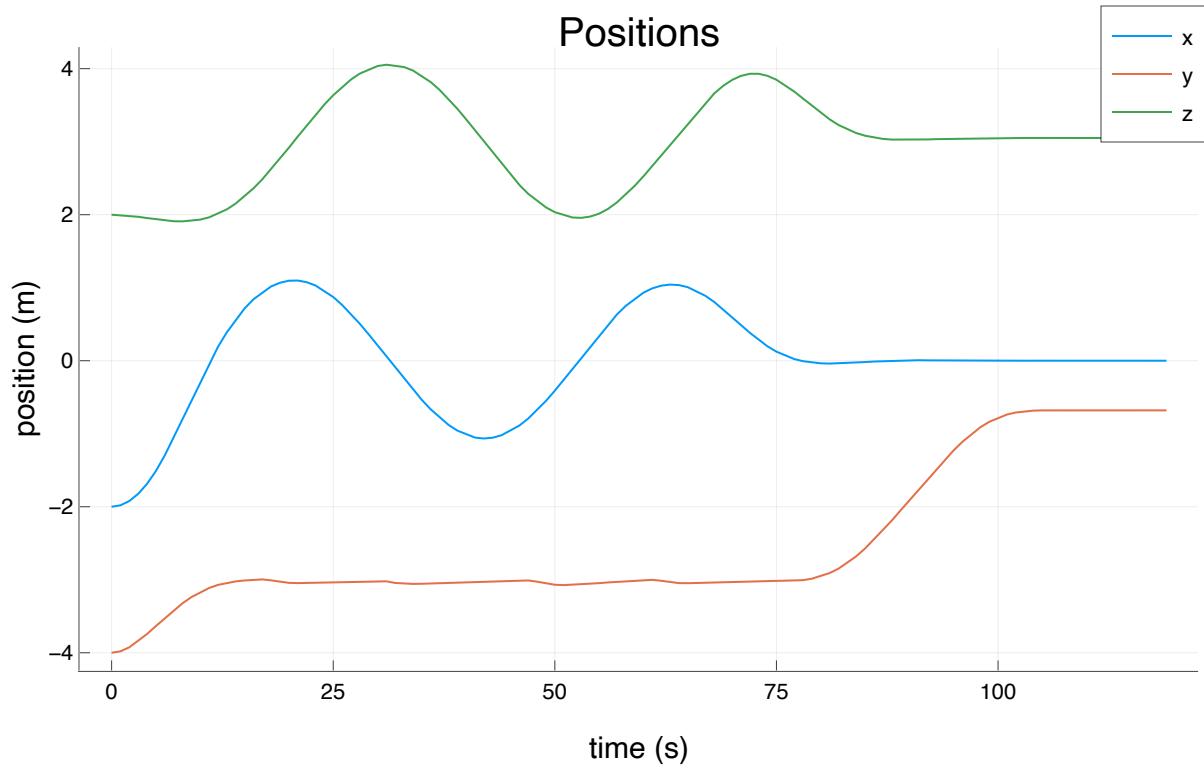
Xm = mat_from_vec(X_sim)
Um = mat_from_vec(U_sim)
display(plot(t_vec,Xm[1:3,:]',title = "Positions",
            xlabel = "time (s)", ylabel = "position (m)",
            label = ["x" "y" "z"]))
display(plot(t_vec,Xm[4:6,:]',title = "Velocities",
            xlabel = "time (s)", ylabel = "velocity (m/s)",
            label = ["x" "y" "z"]))
display(plot(t_vec[1:end-1],Um',title = "Control",
            xlabel = "time (s)", ylabel = "thrust (N)",
            label = ["x" "y" "z"]))

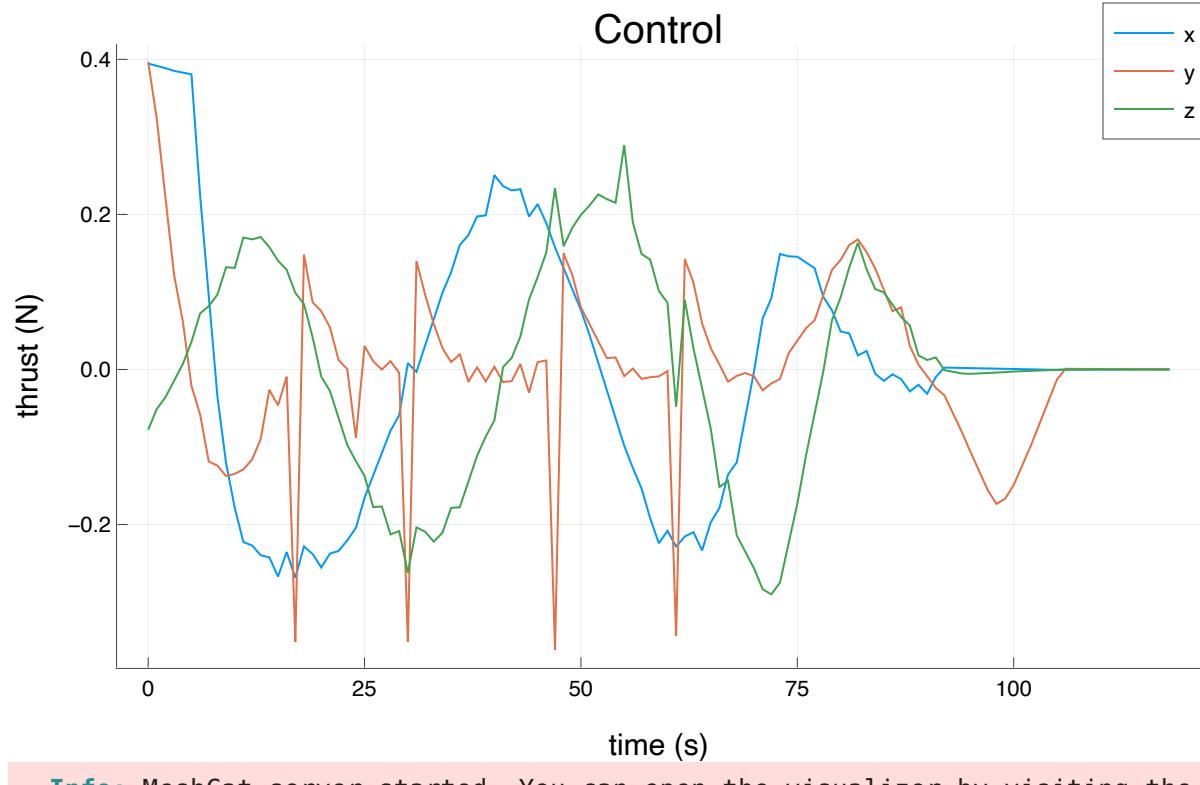
display/animate_rendezvous(X_sim, X_ref, dt;show_reference = false))
# -----plotting/animation-----

# tests
@test norm(X_sim[end] - xg) < 1e-3 # goal
xs=[x[1] for x in X_sim]
ys=[x[2] for x in X_sim]
zs=[x[3] for x in X_sim]
@test maximum(ys) <= (xg[2] + 1e-3)
@test maximum(zs) >= 4 # check to see if you did the circle
@test minimum(zs) <= 2 # check to see if you did the circle
@test maximum(xs) >= 1 # check to see if you did the circle
@test maximum(norm.(U_sim,Inf)) <= 0.4 + 1e-3 # control constraints satisfied

```

```
size(X_ref) = (200,)  
simulating 2%|███████████ | ETA: 0:01:01 ⚠ Warning: Problem status INFEASIBLE; solution may be inaccurate.  
└ @ Convex ~/.julia/packages/Convex/tSTAW/src/solution.jl:342  
simulating 24%|███████████ | ETA: 0:00:06 ⚠ Warning: Problem status INFEASIBLE; solution may be inaccurate.  
└ @ Convex ~/.julia/packages/Convex/tSTAW/src/solution.jl:342  
simulating 39%|███████████ | ETA: 0:00:03 ⚠ Warning: Problem status INFEASIBLE; solution may be inaccurate.  
└ @ Convex ~/.julia/packages/Convex/tSTAW/src/solution.jl:342  
simulating 51%|███████████ | ETA: 0:00:02 ⚠ Warning: Problem status INFEASIBLE; solution may be inaccurate.  
└ @ Convex ~/.julia/packages/Convex/tSTAW/src/solution.jl:342  
simulating 100%|███████████ | Time: 0:00:03
```





Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:
<http://127.0.0.1:8704>

[Open Controls](#)



Test Summary: | Pass Total
 convex mpc | 6 6

Out[7]: `Test.DefaultTestSet("convex mpc", Any[], 6, false, false)`

